

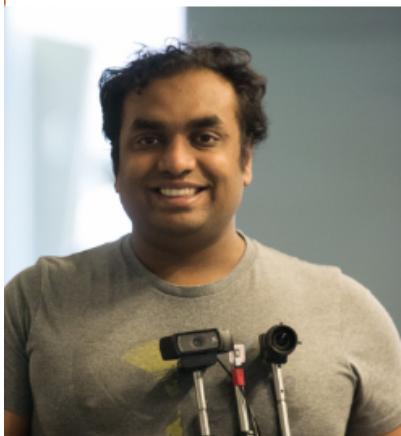
Practical Neural Network Design Using Reinforcement Learning

Bowen Baker

Media Lab

bowen@mit.edu

Co-authors



Otkrist Gupta
MIT Media Lab



Nikhil Naik
Harvard



Ramesh Raskar
MIT Media Lab

Motivation

- Neural network design is still hand crafted.

Motivation

- Neural network design is still hand crafted.
- Despite the wide usage of a few main networks, we may want networks specialized for specific tasks.

Motivation

- Neural network design is still hand crafted.
- Despite the wide usage of a few main networks, we may want networks specialized for specific tasks.
- We may want more than 1 specialized model, e.g. for the ensembling purposes.

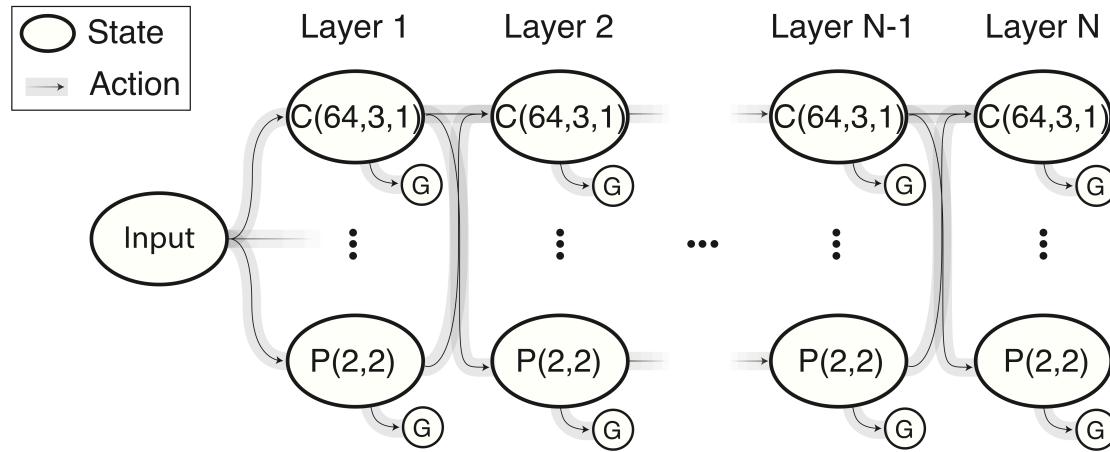
Automating Tasks With Reinforcement Learning



Outline

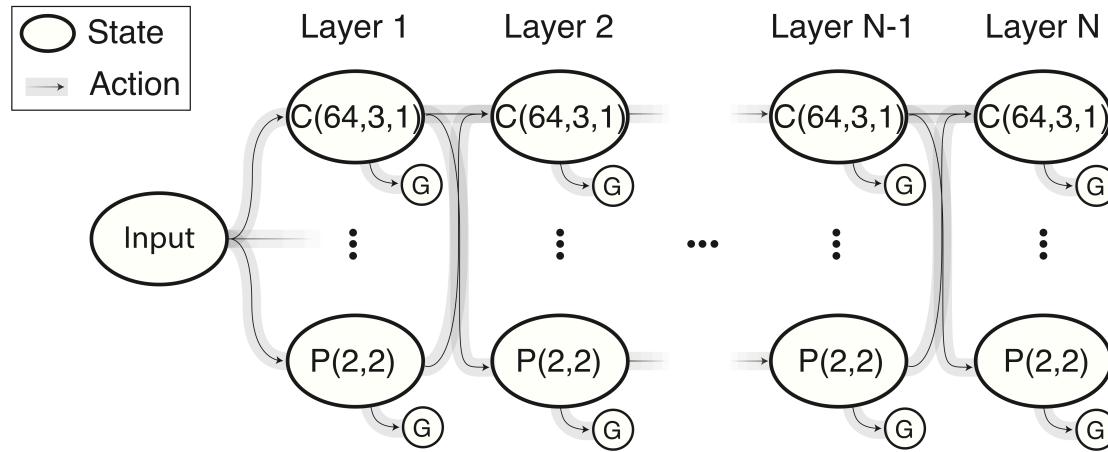
1. Modeling Architecture Selection as a Markov Decision Process
2. Reinforcement Learning Background
3. Results with Q-Learning
4. Accelerating Architecture Selection with Simple Early Stopping Algorithms

Modeling Architecture Selection as a Markov Decision Process



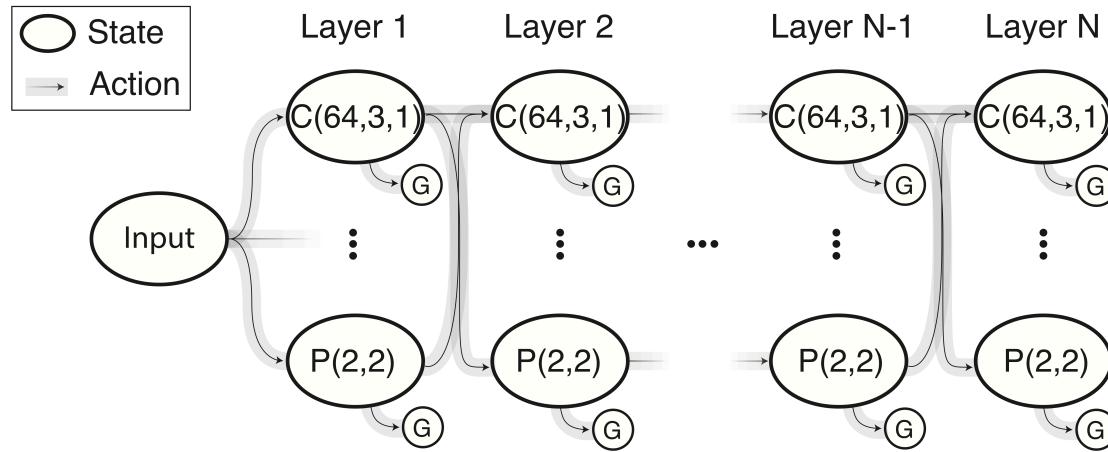
- $C(64,3,1)$ – Convolutional Layer with 64 learnable kernels, 3x3 kernel size, and stride of 1

Modeling Architecture Selection as a Markov Decision Process



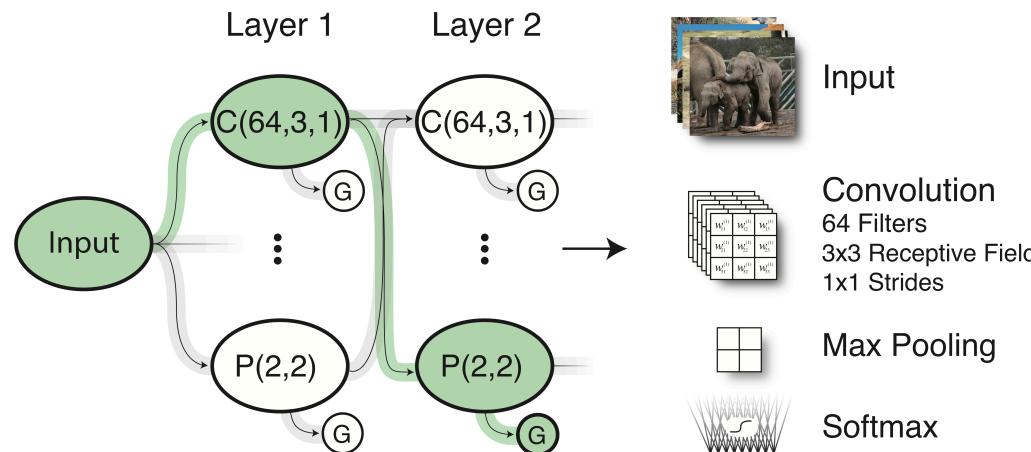
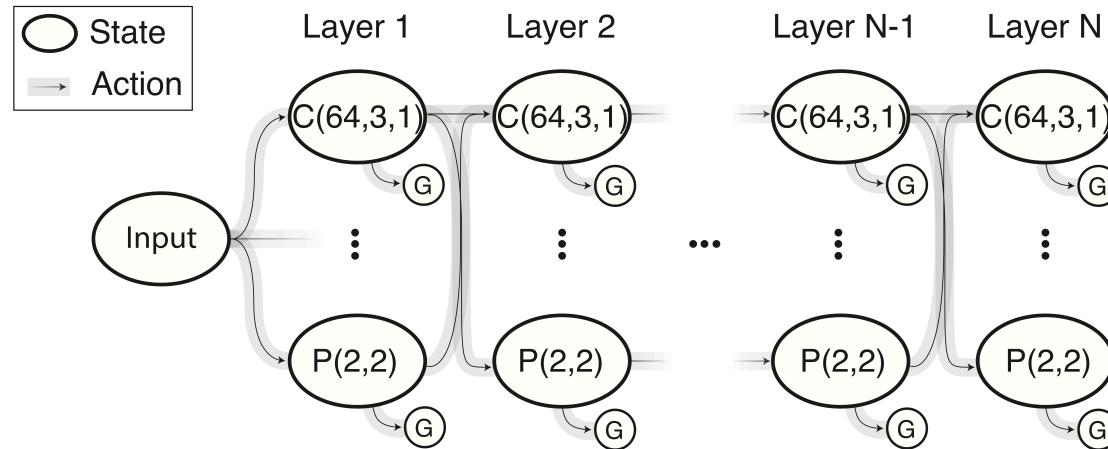
- $C(64,3,1)$ – Convolutional Layer with 64 learnable kernels, 3x3 kernel size, and stride of 1
- $P(2,2)$ – Max Pooling Layer with 2x2 kernel size and stride 2

Modeling Architecture Selection as a Markov Decision Process



- $C(64,3,1)$ – Convolutional Layer with 64 learnable kernels, 3x3 kernel size, and stride of 1
- $P(2,2)$ – Max Pooling Layer with 2x2 kernel size and stride 2
- G – Termination State (e.g. Softmax)

Modeling Architecture Selection as a Markov Decision Process



Q-Learning

$Q^*(s, u)$ -- Denotes the expected reward when following an optimal policy after taking action u at state s

Q-Learning

$$Q^*(s_i, u) = \mathbb{E} \left[r + \gamma \max_{u' \in \mathcal{U}(s_j)} Q^*(s_j, u') \right]$$

γ -- Discount Factor

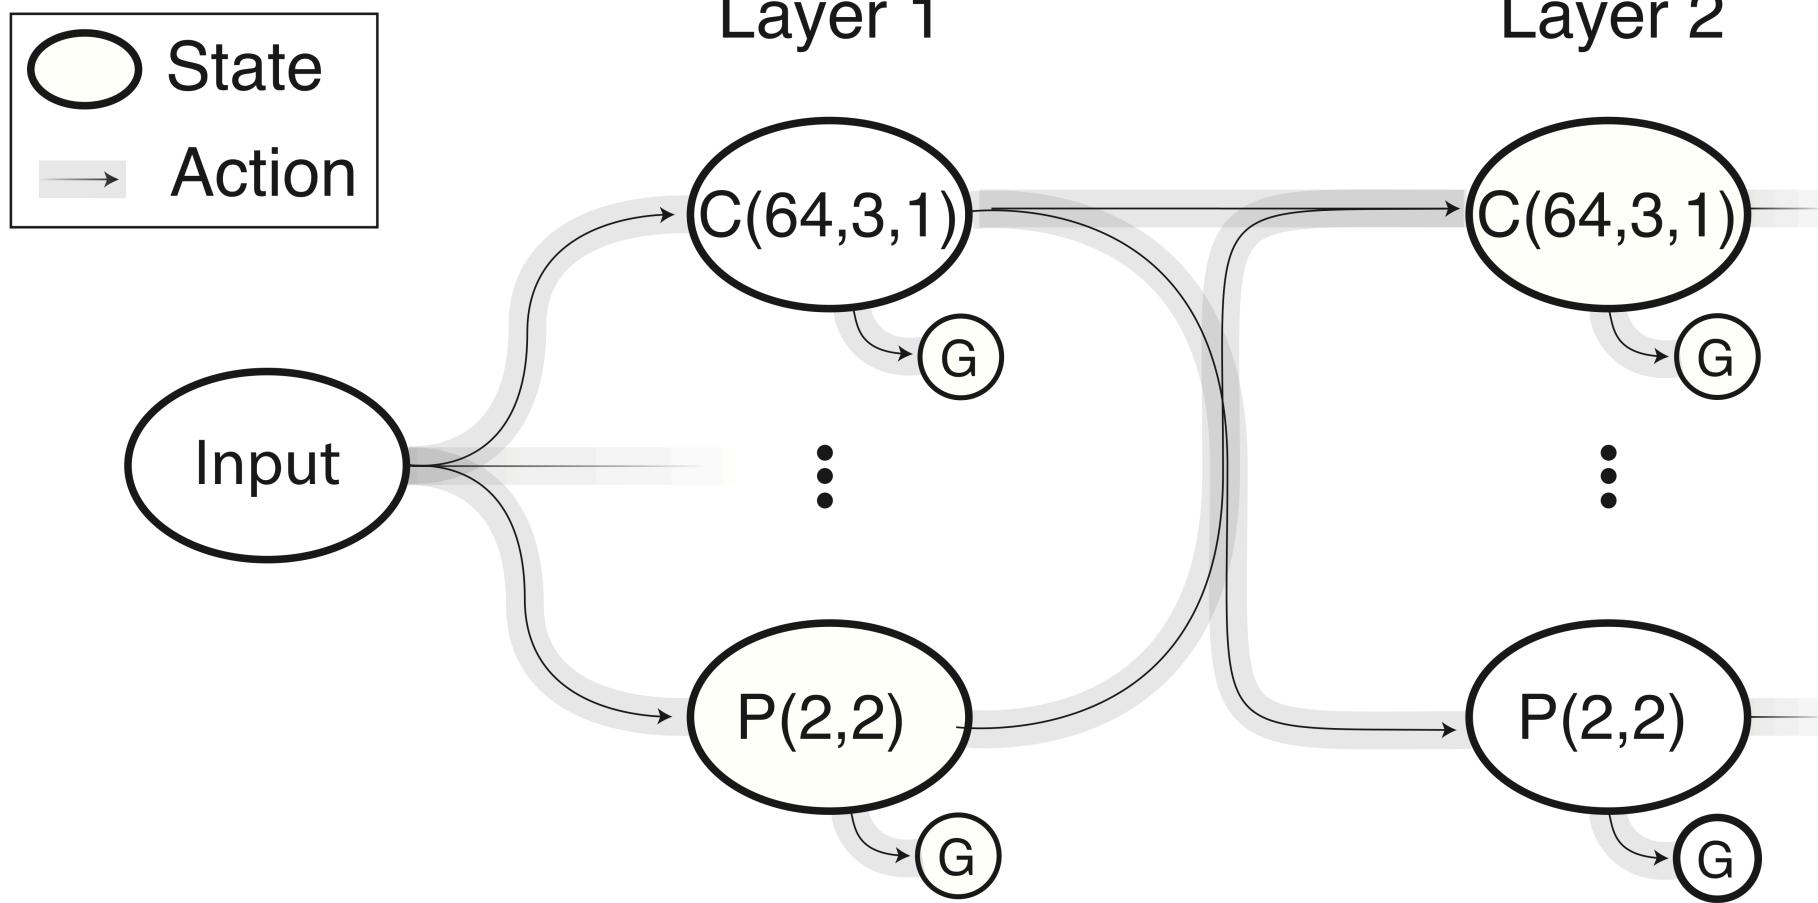
r -- Reward received from
the (s_i, u, s_j) transition

Q-Learning

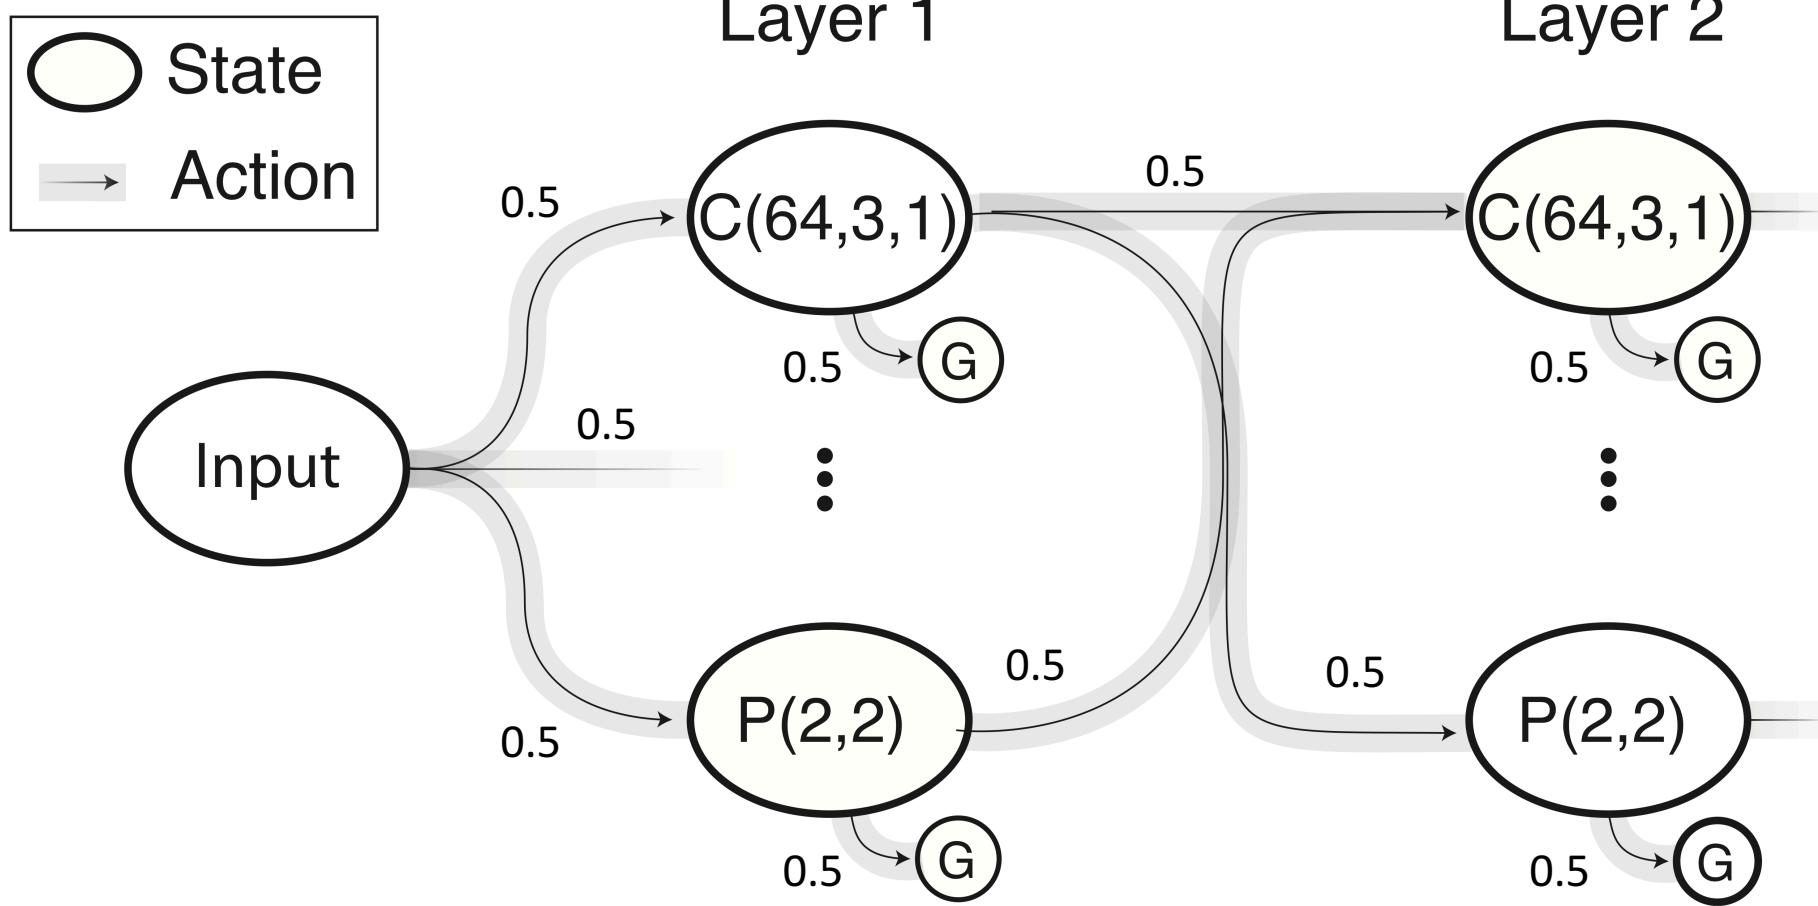
$$Q^*(s_i, u) = \mathbb{E} \left[r + \gamma \max_{u' \in \mathcal{U}(s_j)} Q^*(s_j, u') \right]$$

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right]$$

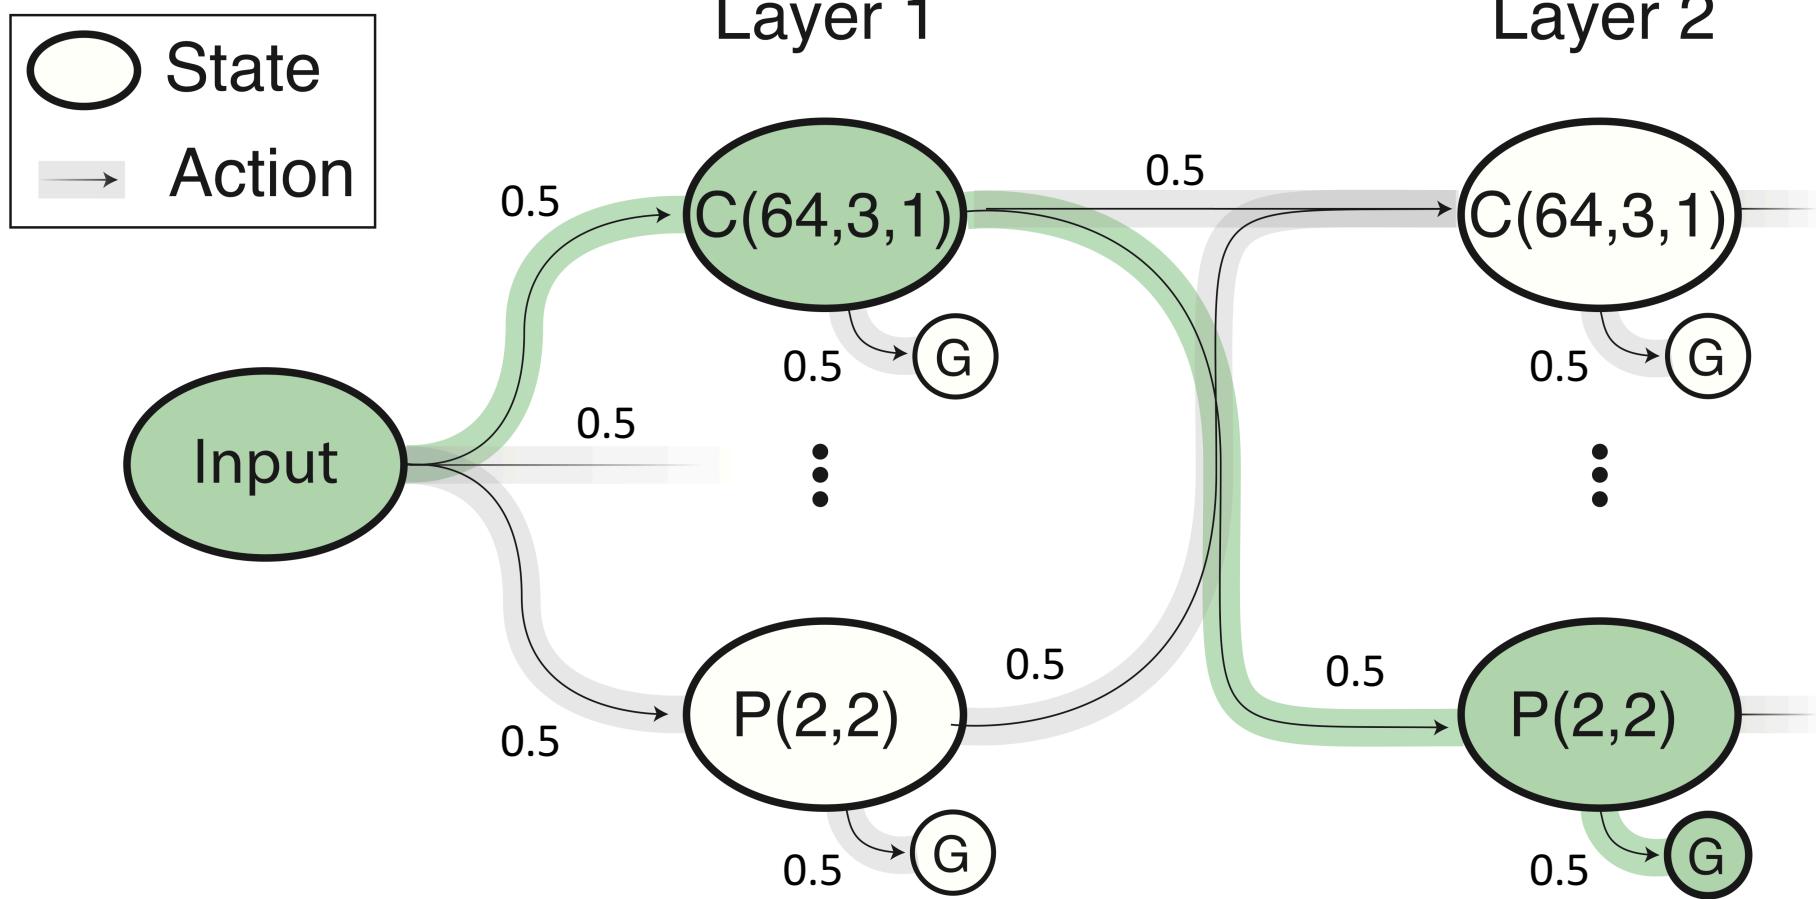
Q-Value Update (Example)



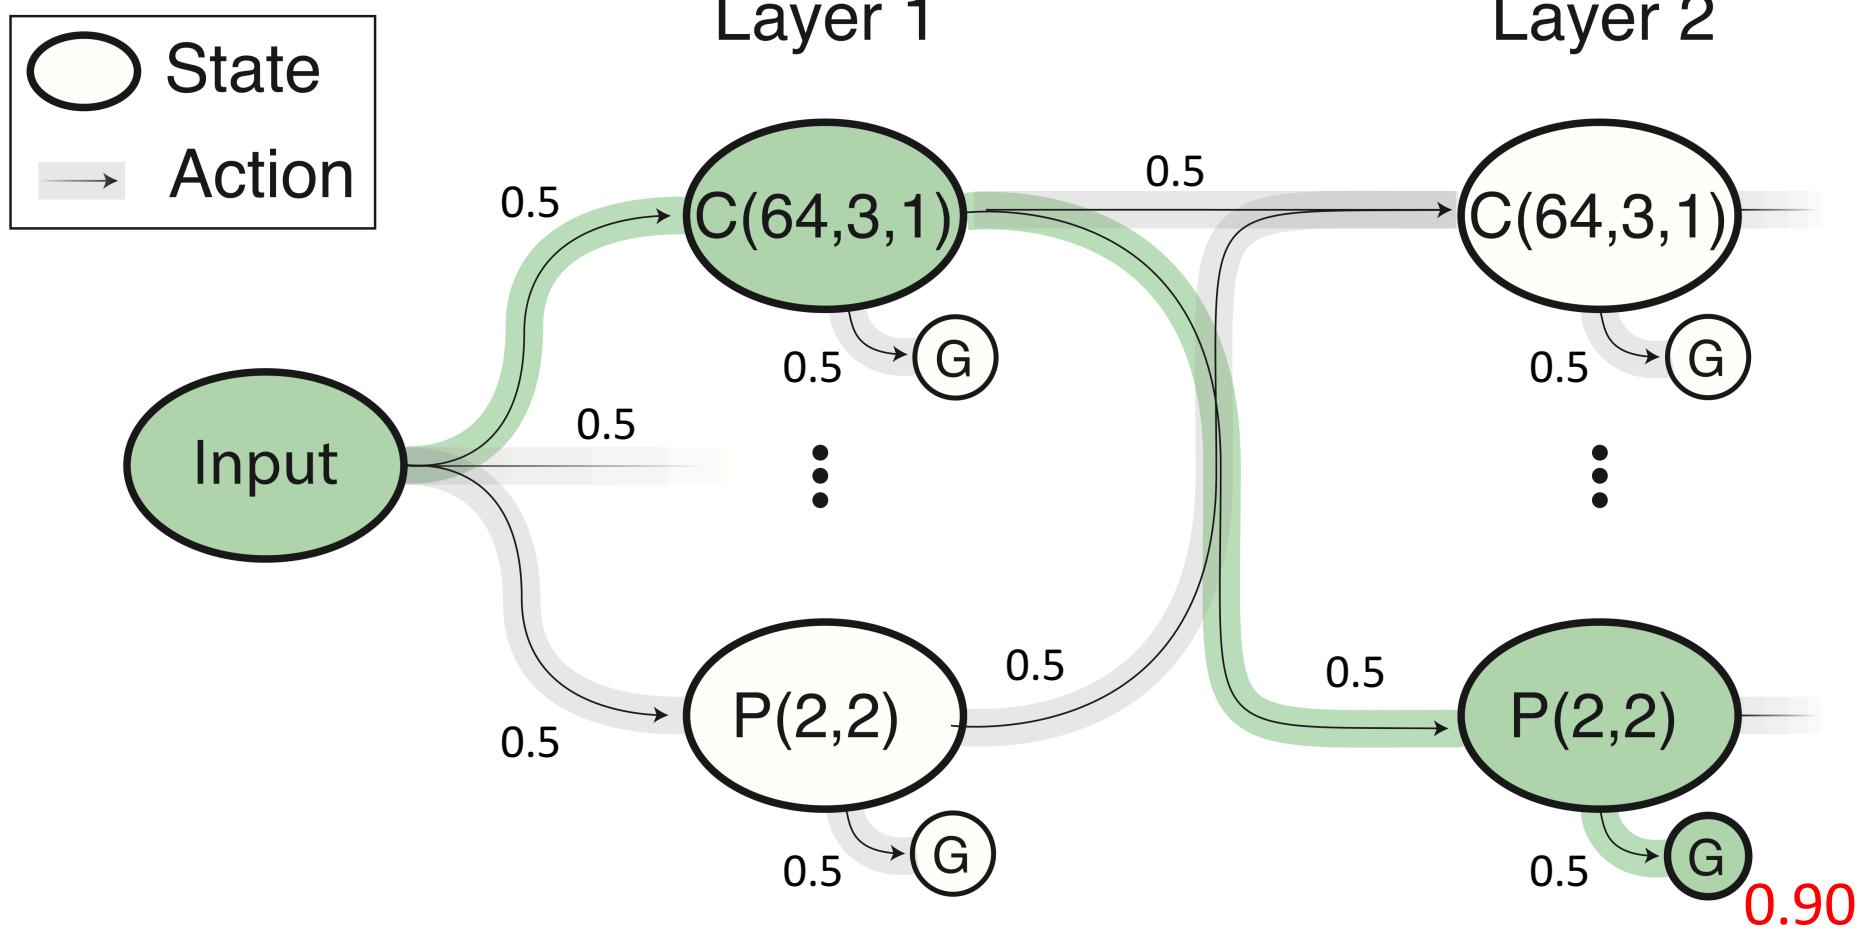
Q-Value Update (Example)



Q-Value Update (Example)

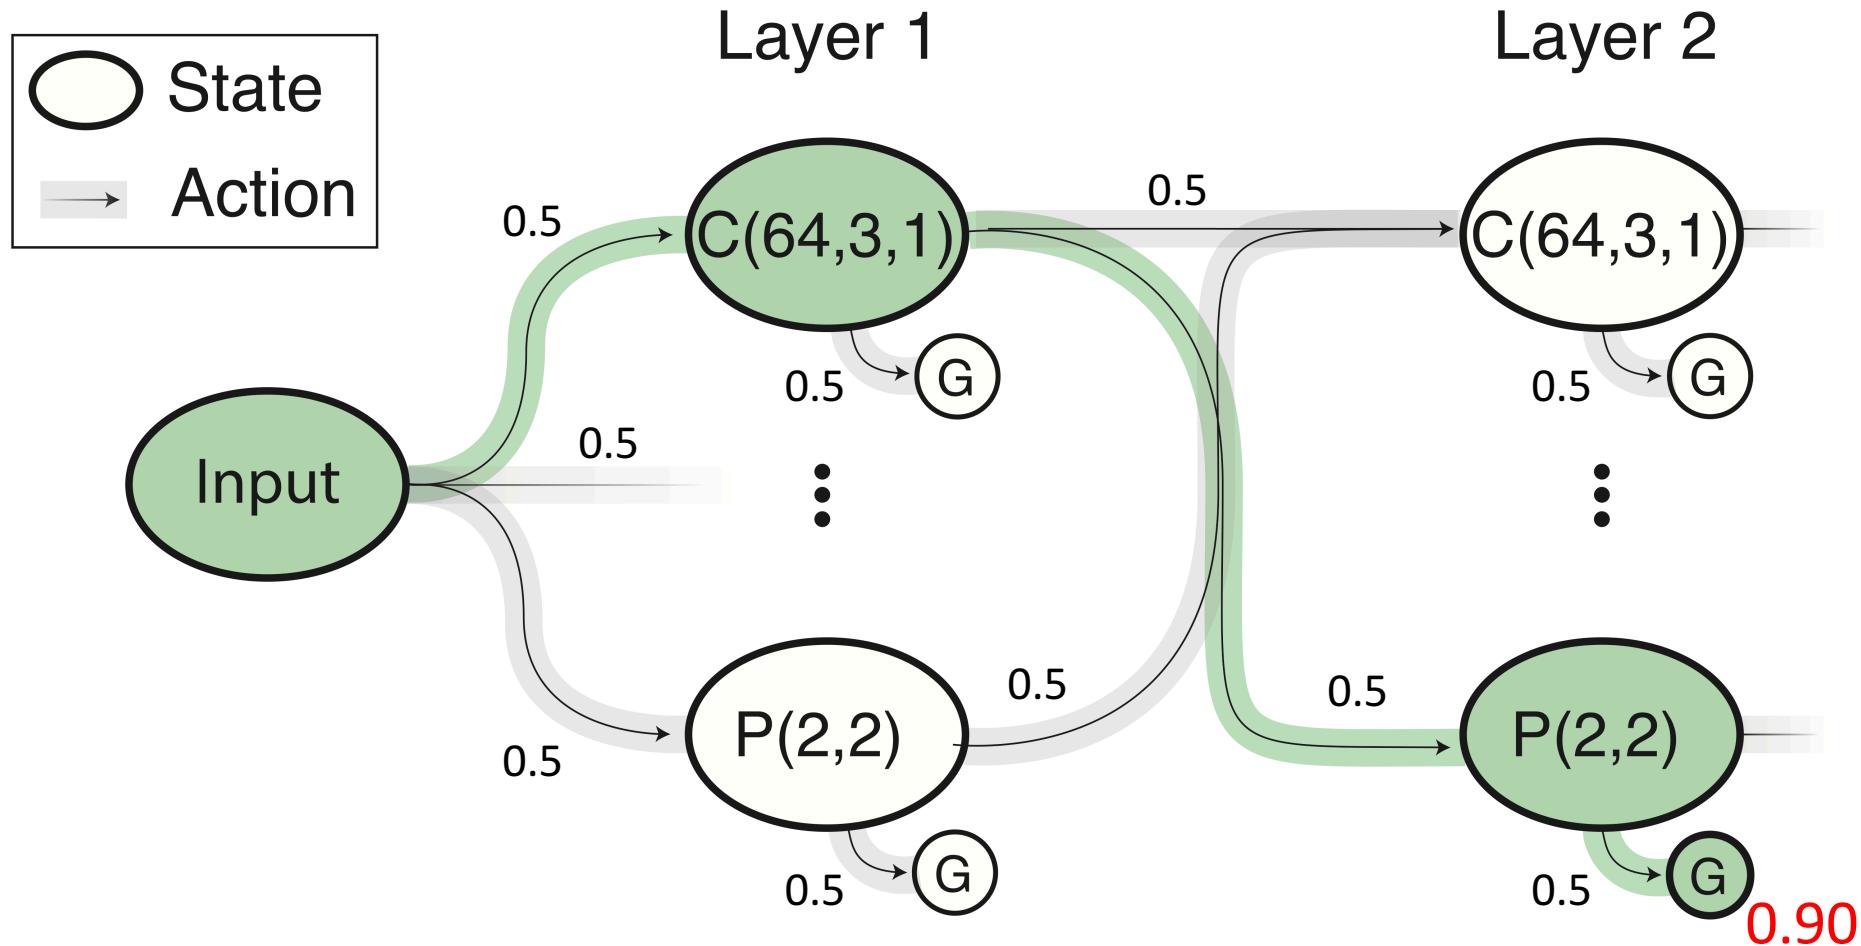


Q-Value Update (Example)



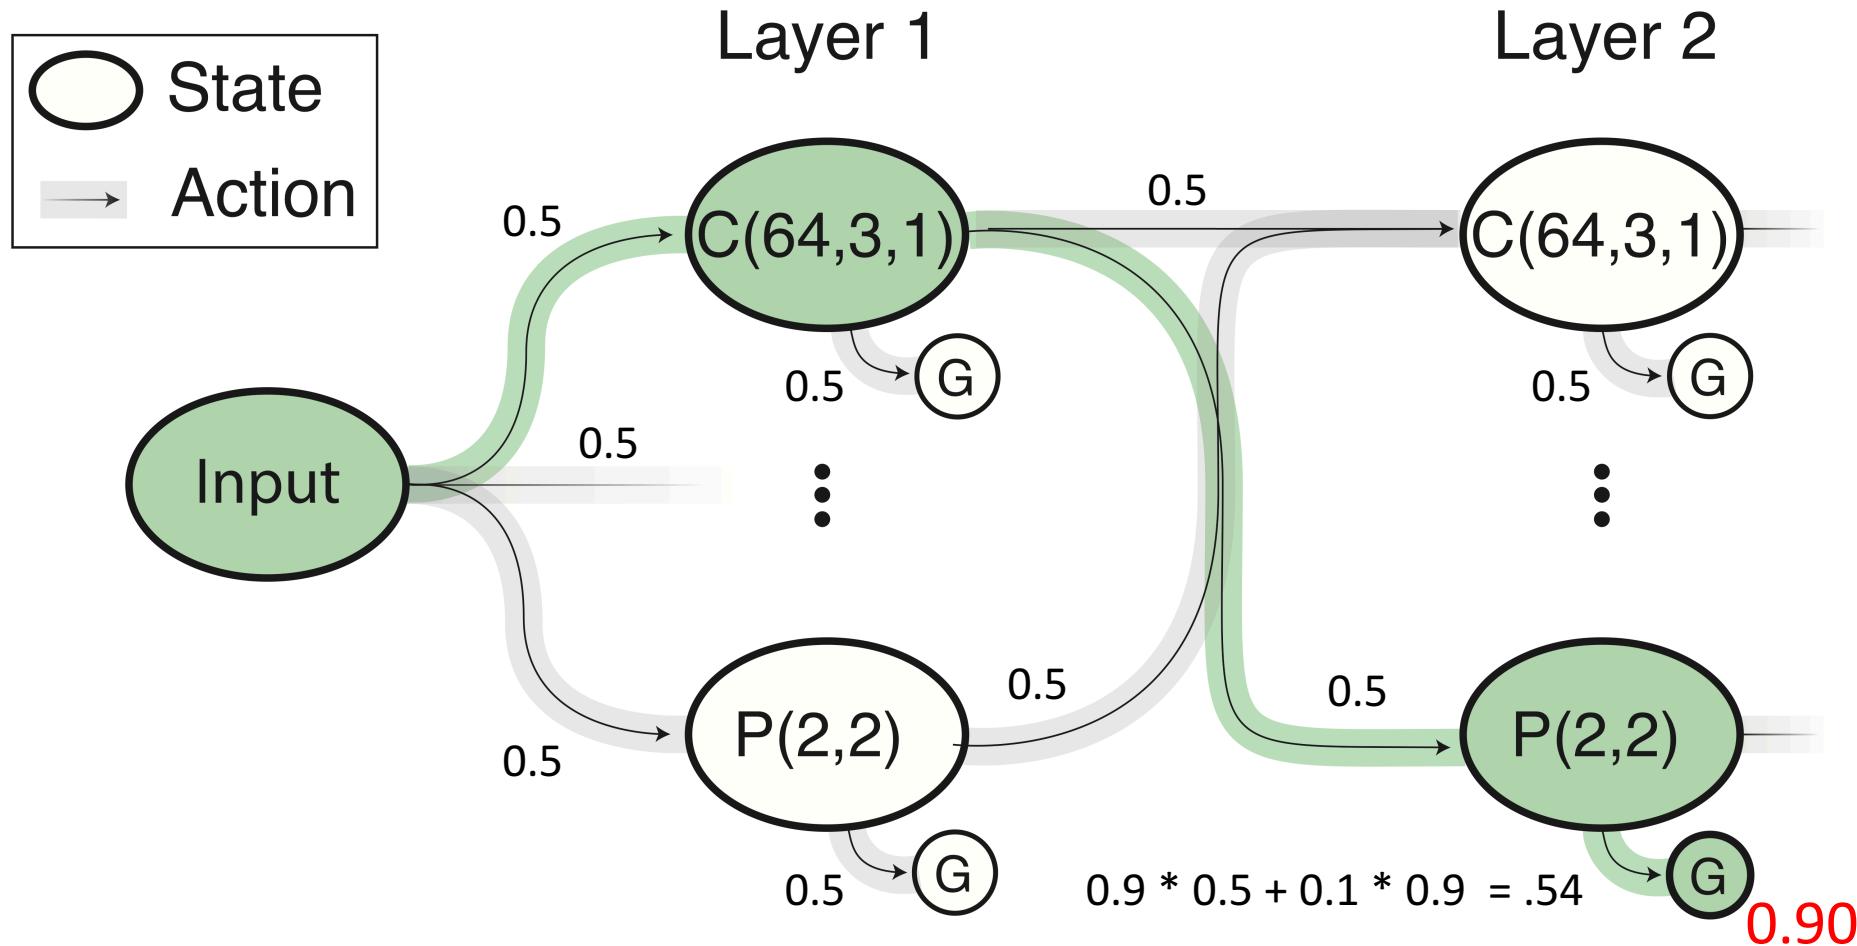
Q-Value Update (Example)

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right], \quad \gamma = 1, \quad \alpha = 0.1$$



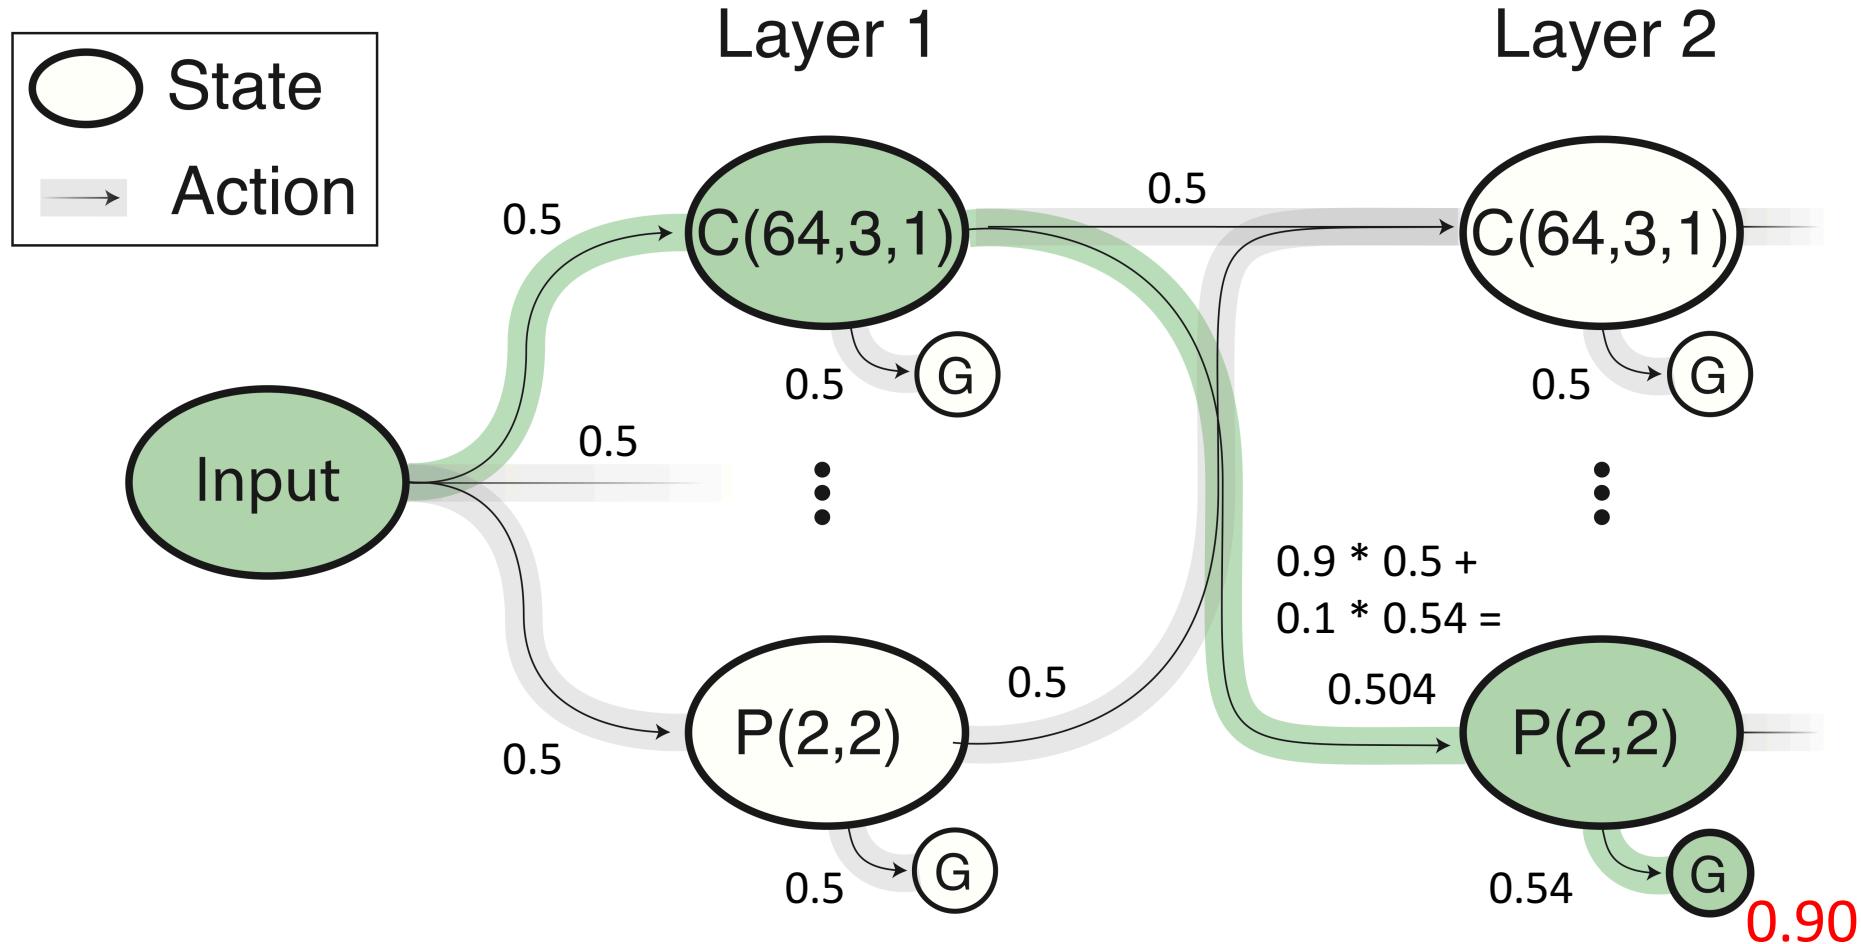
Q-Value Update (Example)

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right], \quad \gamma = 1, \quad \alpha = 0.1$$



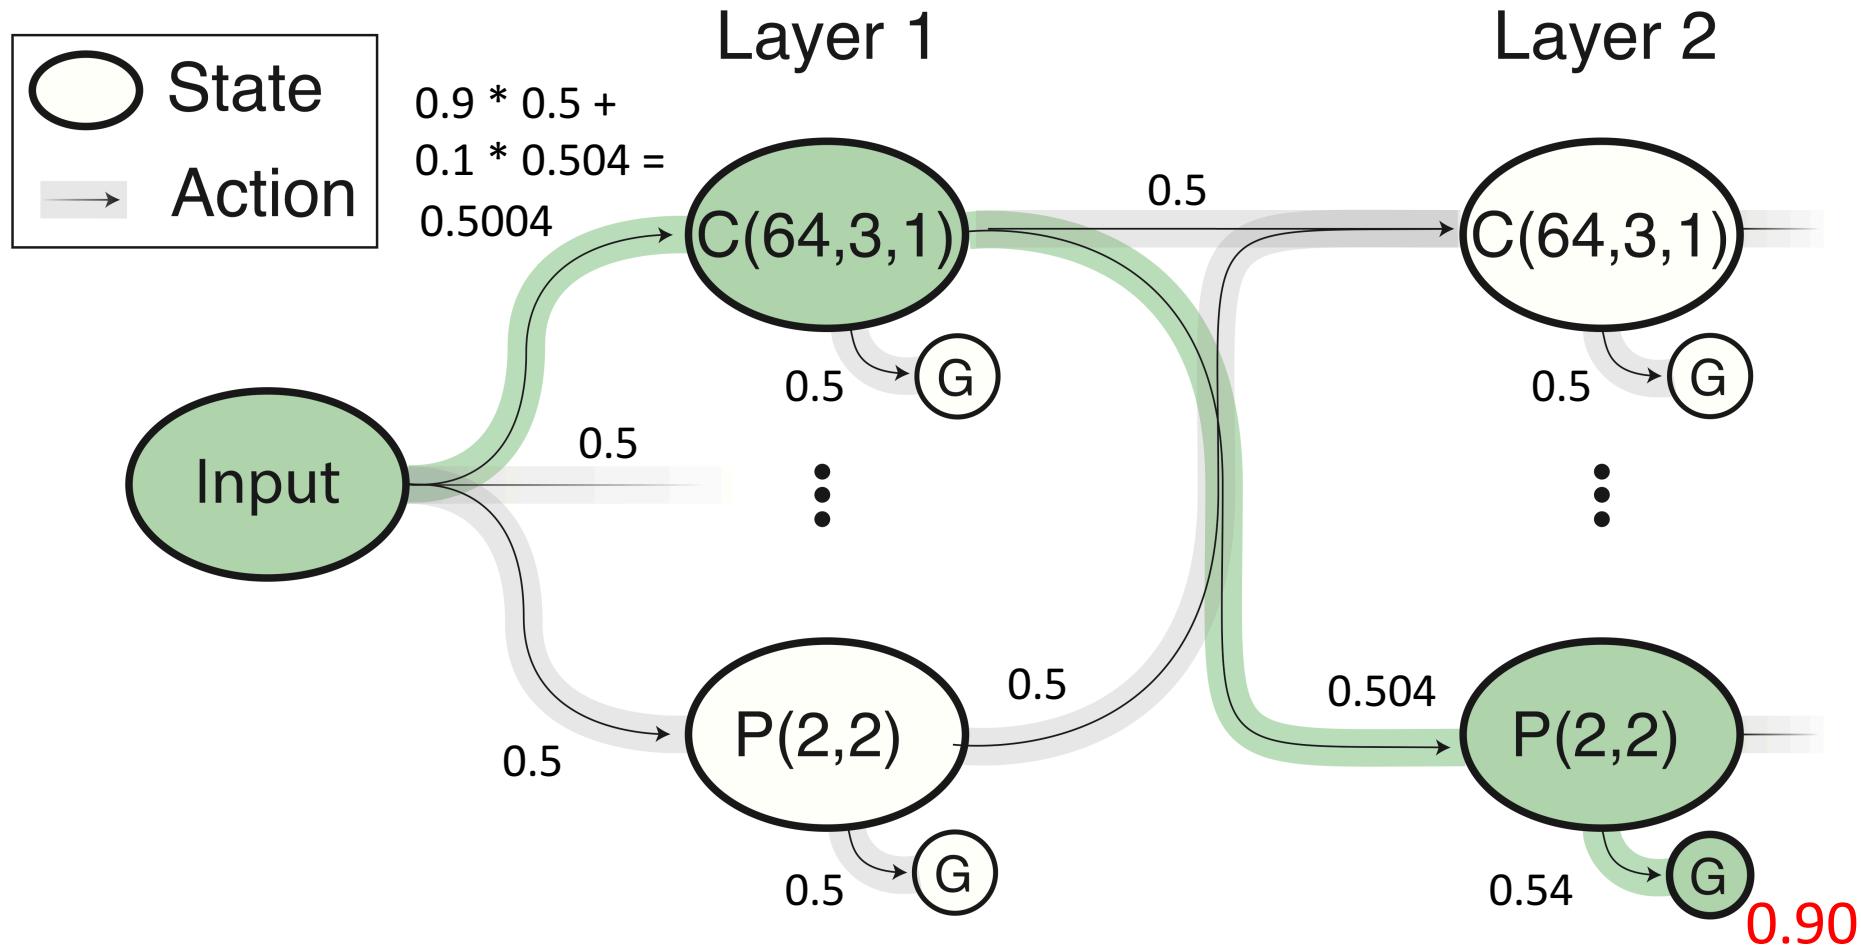
Q-Value Update (Example)

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right], \quad \gamma = 1, \quad \alpha = 0.1$$

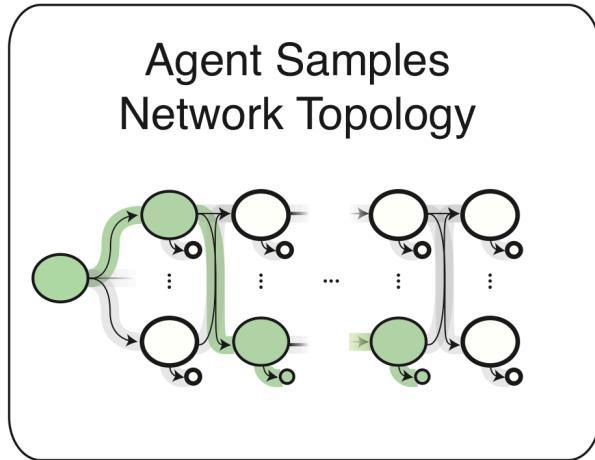


Q-Value Update (Example)

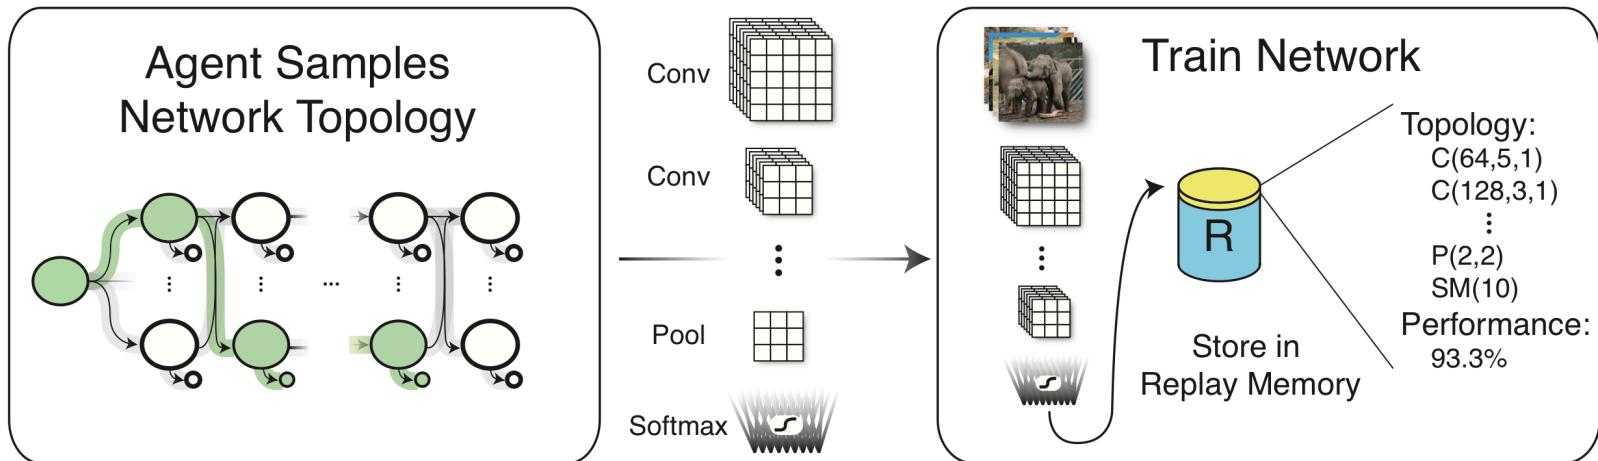
$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right], \quad \gamma = 1, \quad \alpha = 0.1$$



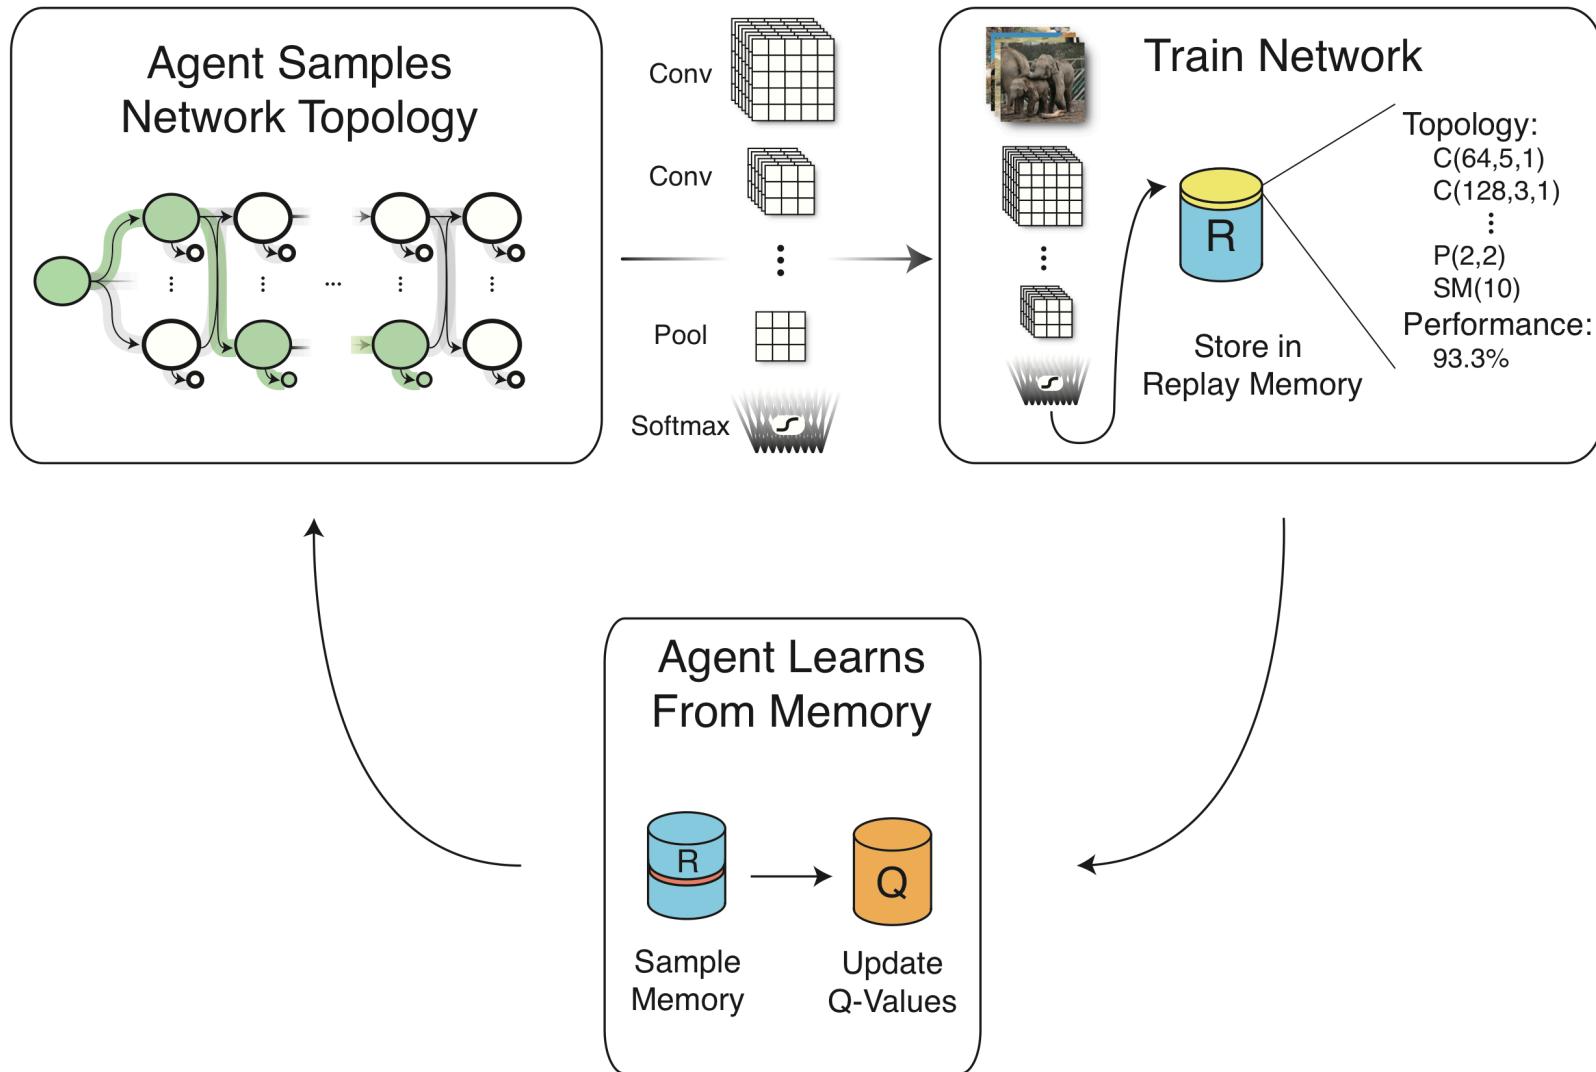
MetaQNN



MetaQNN



MetaQNN



Sampling Networks

Epsilon-Greedy Exploration:

- State s corresponds to the last layer chosen
- Action u corresponds to the next layer chosen

$$u = \begin{cases} \text{Uniform}[\mathcal{U}(s)] & \text{with probability } \epsilon \\ \arg \max_{u' \in \mathcal{U}(s)} [Q(s, u')] & \text{with probability } 1 - \epsilon \end{cases}$$

Sampling Networks

Epsilon-Greedy Exploration:

- State s corresponds to the last layer chosen
- Action u corresponds to the next layer chosen

$$u = \begin{cases} \text{Uniform}[\mathcal{U}(s)] & \text{with probability } \epsilon \\ \arg \max_{u' \in \mathcal{U}(s)} [Q(s, u')] & \text{with probability } 1 - \epsilon \end{cases}$$

| ϵ | 1.0 | 0.9 | 0.8 | 0.7 | 0.6 | 0.5 | 0.4 | 0.3 | 0.2 | 0.1 |
|------------------|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| # Models Trained | 1500 | 100 | 100 | 100 | 150 | 150 | 150 | 150 | 150 | 150 |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|---|--|
| Convolution (C) | $i \sim \text{Layer depth}$ $f \sim \text{Receptive field size}$ $\ell \sim \text{Stride}$ $d \sim \# \text{ receptive fields}$ $n \sim \text{Representation size}$ | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim \text{Layer depth}$ $(f, \ell) \sim (\text{Receptive field size}, \text{Strides})$ $n \sim \text{Representation size}$ | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim \text{Layer depth}$ $n \sim \# \text{ consecutive FC layers}$ $d \sim \# \text{ neurons}$ | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim \text{Previous State}$ $t \sim \text{Type}$ | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|---|--|
| Convolution (C) | $i \sim \text{Layer depth}$ $f \sim \text{Receptive field size}$ $\ell \sim \text{Stride}$ $d \sim \# \text{ receptive fields}$ $n \sim \text{Representation size}$ | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim \text{Layer depth}$ $(f, \ell) \sim (\text{Receptive field size}, \text{Strides})$ $n \sim \text{Representation size}$ | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim \text{Layer depth}$ $n \sim \# \text{ consecutive FC layers}$ $d \sim \# \text{ neurons}$ | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim \text{Previous State}$ $t \sim \text{Type}$ | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

State Space

| Layer Type | Layer Parameters | Parameter Values |
|----------------------|--|--|
| Convolution (C) | $i \sim$ Layer depth $f \sim$ Receptive field size $\ell \sim$ Stride $d \sim$ # receptive fields $n \sim$ Representation size | < 12 Square. $\in \{1, 3, 5\}$ Square. Always equal to 1 $\in \{64, 128, 256, 512\}$ $\in \{(\infty, 8], (8, 4], (4, 1]\}$ |
| Pooling (P) | $i \sim$ Layer depth $(f, \ell) \sim$ (Receptive field size, Strides) $n \sim$ Representation size | < 12 Square. $\in \{(5, 3), (3, 2), (2, 2)\}$ $\in \{(\infty, 8], (8, 4] \text{ and } (4, 1]\}$ |
| Fully Connected (FC) | $i \sim$ Layer depth $n \sim$ # consecutive FC layers $d \sim$ # neurons | < 12 < 3 $\in \{512, 256, 128\}$ |
| Termination State | $s \sim$ Previous State $t \sim$ Type | Global Avg. Pooling/Softmax |

Action Space

- Convolution → Any Other Layer

Action Space

- Convolution → Any Other Layer
- Pooling → Any Other Layer / Pooling

Action Space

- Convolution → Any Other Layer
- Pooling → Any Other Layer / Pooling
- Any Layer → Fully Connected
 - if representation size less than 8

Action Space

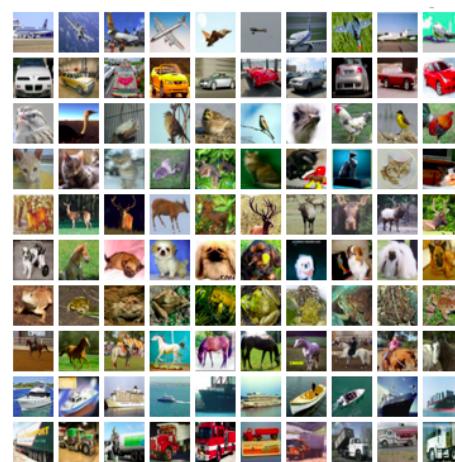
- Convolution → Any Other Layer
- Pooling → Any Other Layer / Pooling
- Any Layer → Fully Connected
 - if representation size less than 8
- Any Layer → Termination

Experiments

MNIST



CIFAR-10



SVHN



- Hand Written Digits
- 60,000 Training Examples
- 10,000 Testing Examples
- 10 classes

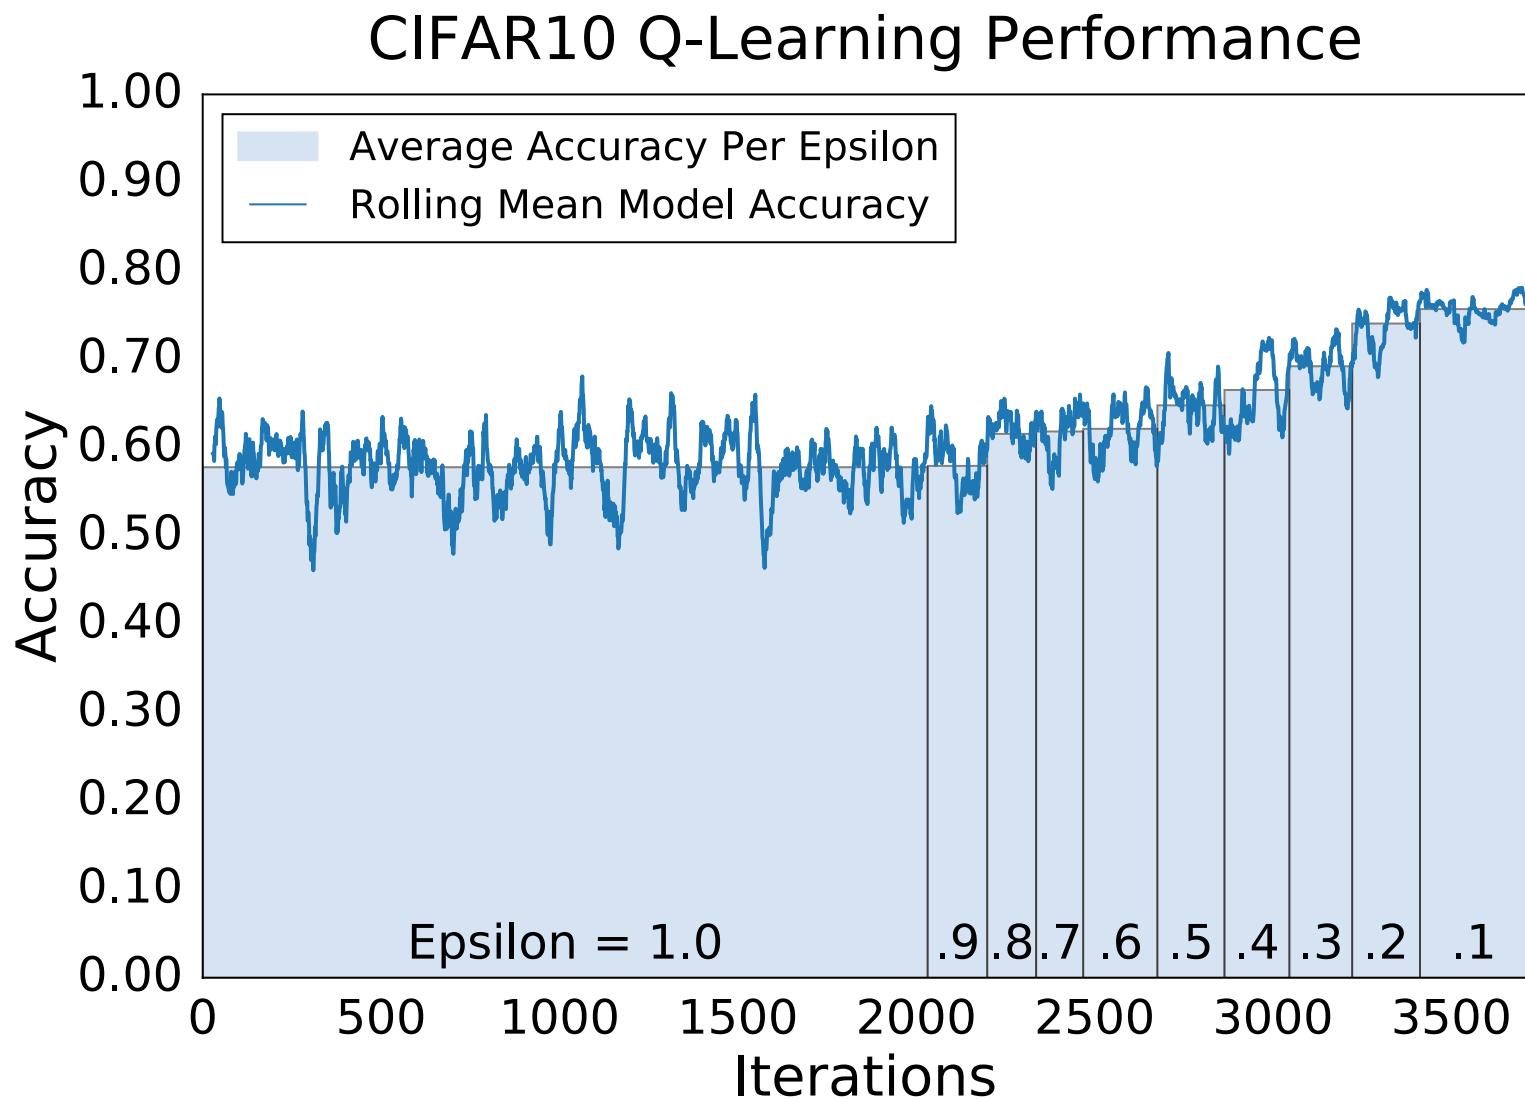
- Tiny Images
- 50,000 Training Examples
- 10,000 Testing Examples
- 10 classes

- Street View House Digits
- 73257 Training Examples
- 26032 Testing Examples
- 531131 ‘Extra’ Examples
- 10 classes

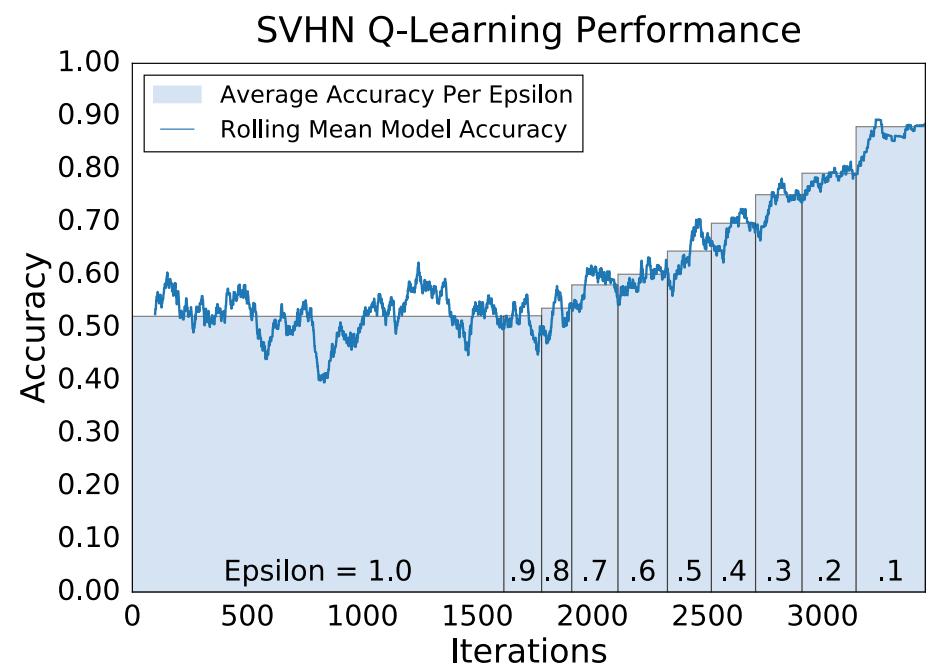
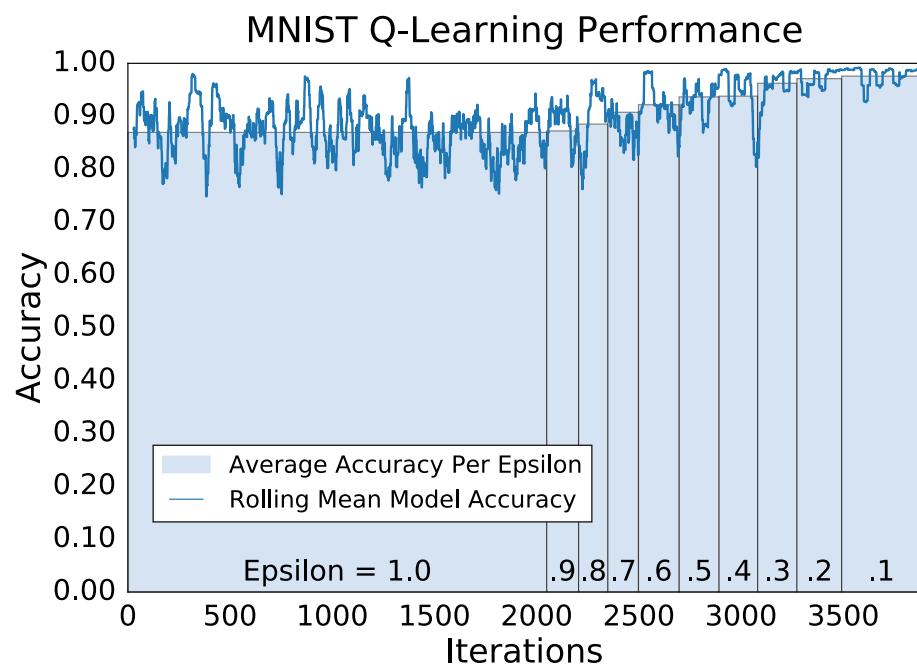
Hardware

- ~10 GPU's
 - Mostly 2015 Titan Xs
 - Some GTX 1080s
- Each experiment took ~10 days
 - Roughly 100 GPUdays

Results



Results



Results

Comparison Against Models with similar design modules:

| Method | CIFAR-10 | SVHN | MNIST | CIFAR-100 |
|-------------------------------------|-------------|-------------|-------------|---------------|
| Maxout (Goodfellow et al., 2013) | 9.38 | 2.47 | 0.45 | 38.57 |
| NIN (Lin et al., 2013) | 8.81 | 2.35 | 0.47 | 35.68 |
| FitNet (Romero et al., 2014) | 8.39 | 2.42 | 0.51 | 35.04 |
| HighWay (Srivastava et al., 2015) | 7.72 | - | - | - |
| VGGnet (Simonyan & Zisserman, 2014) | 7.25 | - | - | - |
| All-CNN (Springenberg et al., 2014) | 7.25 | - | - | 33.71 |
| MetaQNN (ensemble) | 7.32 | 2.06 | 0.32 | - |
| MetaQNN (top model) | 6.92 | 2.28 | 0.44 | 27.14* |

Results

Comparison Against Models with similar design modules:

| Method | CIFAR-10 | SVHN | MNIST | CIFAR-100 |
|-------------------------------------|-------------|-------------|-------------|---------------|
| Maxout (Goodfellow et al., 2013) | 9.38 | 2.47 | 0.45 | 38.57 |
| NIN (Lin et al., 2013) | 8.81 | 2.35 | 0.47 | 35.68 |
| FitNet (Romero et al., 2014) | 8.39 | 2.42 | 0.51 | 35.04 |
| HighWay (Srivastava et al., 2015) | 7.72 | - | - | - |
| VGGnet (Simonyan & Zisserman, 2014) | 7.25 | - | - | - |
| All-CNN (Springenberg et al., 2014) | 7.25 | - | - | 33.71 |
| MetaQNN (ensemble) | 7.32 | 2.06 | 0.32 | - |
| MetaQNN (top model) | 6.92 | 2.28 | 0.44 | 27.14* |

Comparison Against more complex modules:

| Method | CIFAR-10 | SVHN | MNIST | CIFAR-100 |
|---------------------------------|-------------|-------------|-------------|--------------|
| DropConnect (Wan et al., 2013) | 9.32 | 1.94 | 0.57 | - |
| DSN (Lee et al., 2015) | 8.22 | 1.92 | 0.39 | 34.57 |
| R-CNN (Liang & Hu, 2015) | 7.72 | 1.77 | 0.31 | 31.75 |
| MetaQNN (ensemble) | 7.32 | 2.06 | 0.32 | - |
| MetaQNN (top model) | 6.92 | 2.28 | 0.44 | 27.14* |
| Resnet(110) (He et al., 2015) | 6.61 | - | - | - |
| Resnet(1001) (He et al., 2016) | 4.62 | - | - | 22.71 |
| ELU (Clevert et al., 2015) | 6.55 | - | - | 24.28 |
| Tree+Max-Avg (Lee et al., 2016) | 6.05 | 1.69 | 0.31 | 32.37 |

Meta-Modeling Comparison on CIFAR-10

| Method | Test Error on CIFAR-10 | # Samples | Estimated Computation (GPU-Days) |
|--|------------------------|-----------|----------------------------------|
| MetaQNN (Ours) | 6.92 | 2,700 | 100 |
| Neural Architecture Search (Zoph et al., 2016) | 3.65 | 12,800 | 10,000 |
| Large Scale Evolution (Real et al., 2017) | 5.4 | - | 2,600 |
| Bayesian Optimization (Snoek et al., 2012) | 9.5 | 50 | - |

Updated Results:

Different Model Depths Don't Train Equally

| Model Depth | 20 Epoch Accuracy | 300 Epoch Accuracy |
|-------------|-------------------|--------------------|
| 9 | 84.78 | 93.08 |
| 15 | 81.2 | 94.7 |

Updated Results: Different Model Depths Don't Train Equally

| Model Depth | 20 Epoch Accuracy | 300 Epoch Accuracy |
|-------------|-------------------|--------------------|
| 9 | 84.78 | 93.08 |
| 15 | 81.2 | 94.7 |

| Method | CIFAR-10 |
|---------------------------------|---------------------|
| DropConnect (Wan et al., 2013) | 9.32 |
| DSN (Lee et al., 2015) | 8.22 |
| R-CNN (Liang & Hu, 2015) | 7.72 |
| MetaQNN (ensemble) | 7.32 |
| MetaQNN (top model) | 6.92 5.3 |
| Resnet(110) (He et al., 2015) | 6.61 |
| Resnet(1001) (He et al., 2016) | 4.62 |
| ELU (Clevert et al., 2015) | 6.55 |
| Tree+Max-Avg (Lee et al., 2016) | 6.05 |

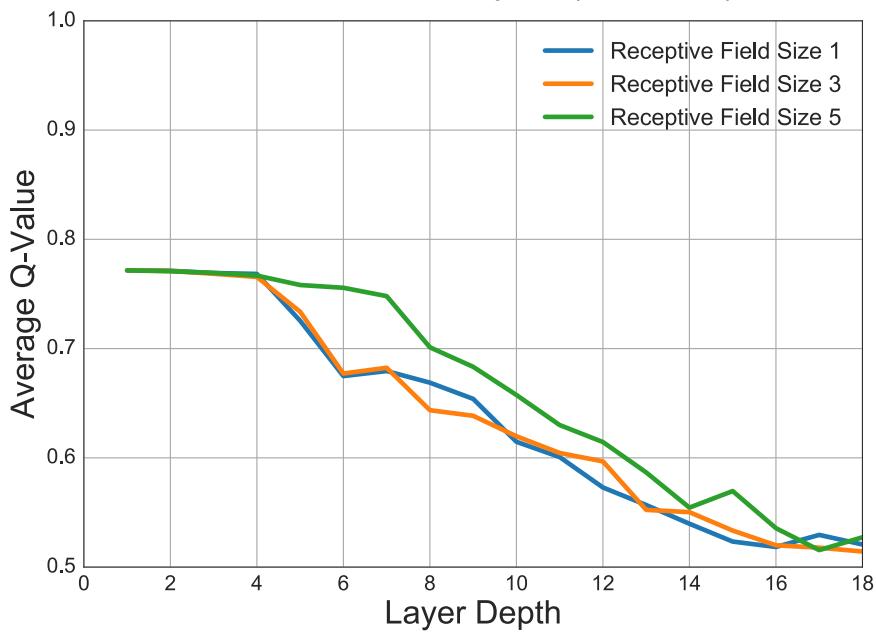
Updated Results:

Different Model Depths Don't Train Equally

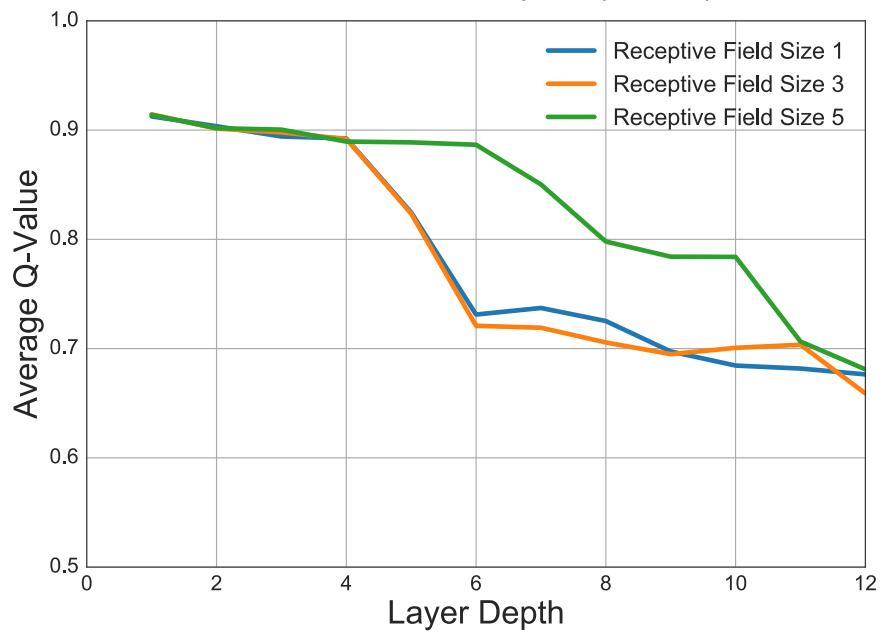
| Method | Test Error on CIFAR-10 | # Samples | Estimated Computation (GPU-Days) |
|--|------------------------|-----------|----------------------------------|
| MetaQNN (Ours) | 5.3 | 2,700 | 100 |
| Neural Architecture Search (Zoph et al., 2016) | 3.65 | 12,800 | 10,000 |
| Large Scale Evolution (Real et al., 2017) | 5.4 | - | 2,600 |
| Bayesian Optimization (Snoek et al., 2012) | 9.5 | 50 | - |

Q-Value Analysis

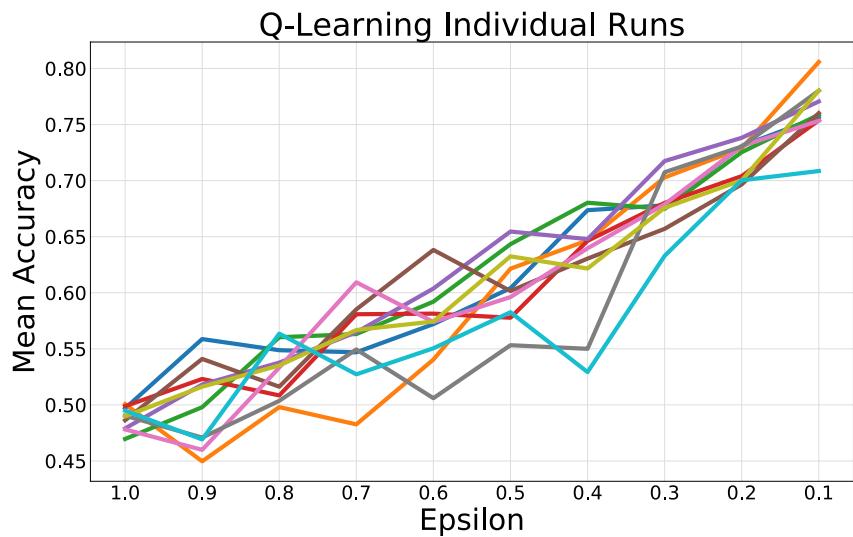
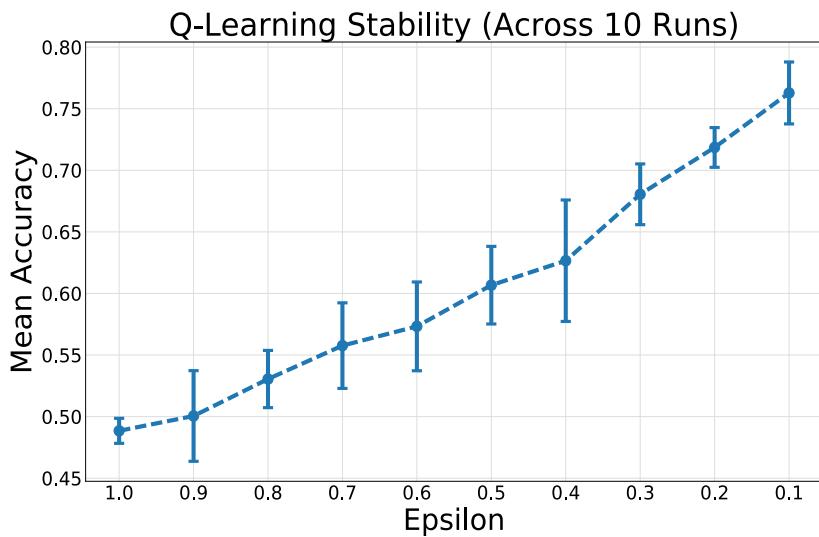
Average Q-Value vs. Layer Depth
for Convolution Layers (CIFAR10)



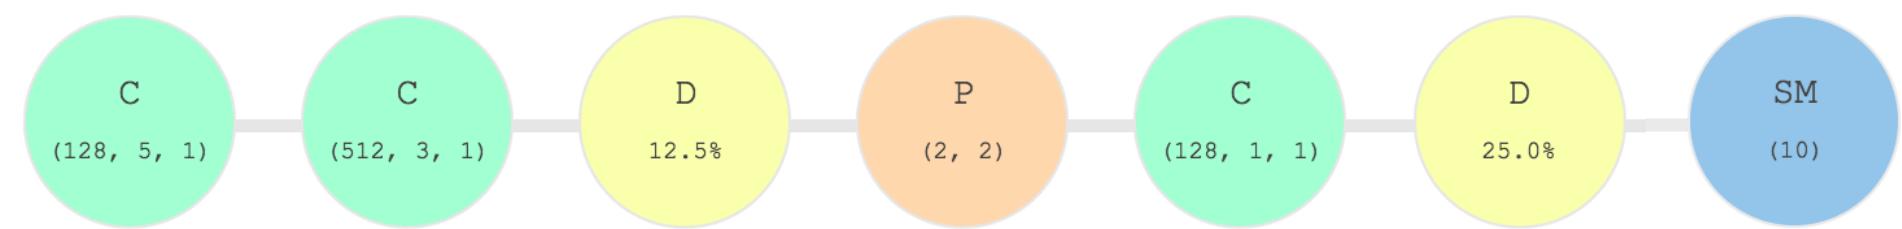
Average Q-Value vs. Layer Depth
for Convolution Layers (SVHN)



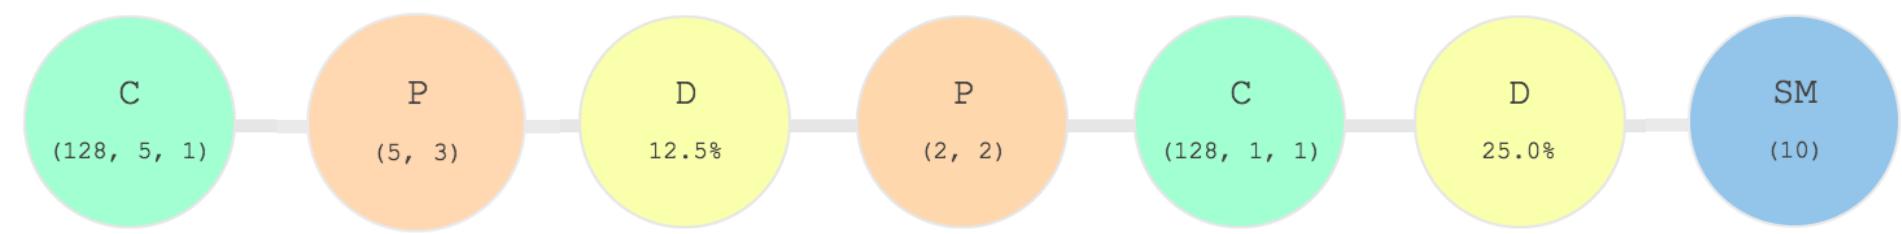
MetaQNN Stability



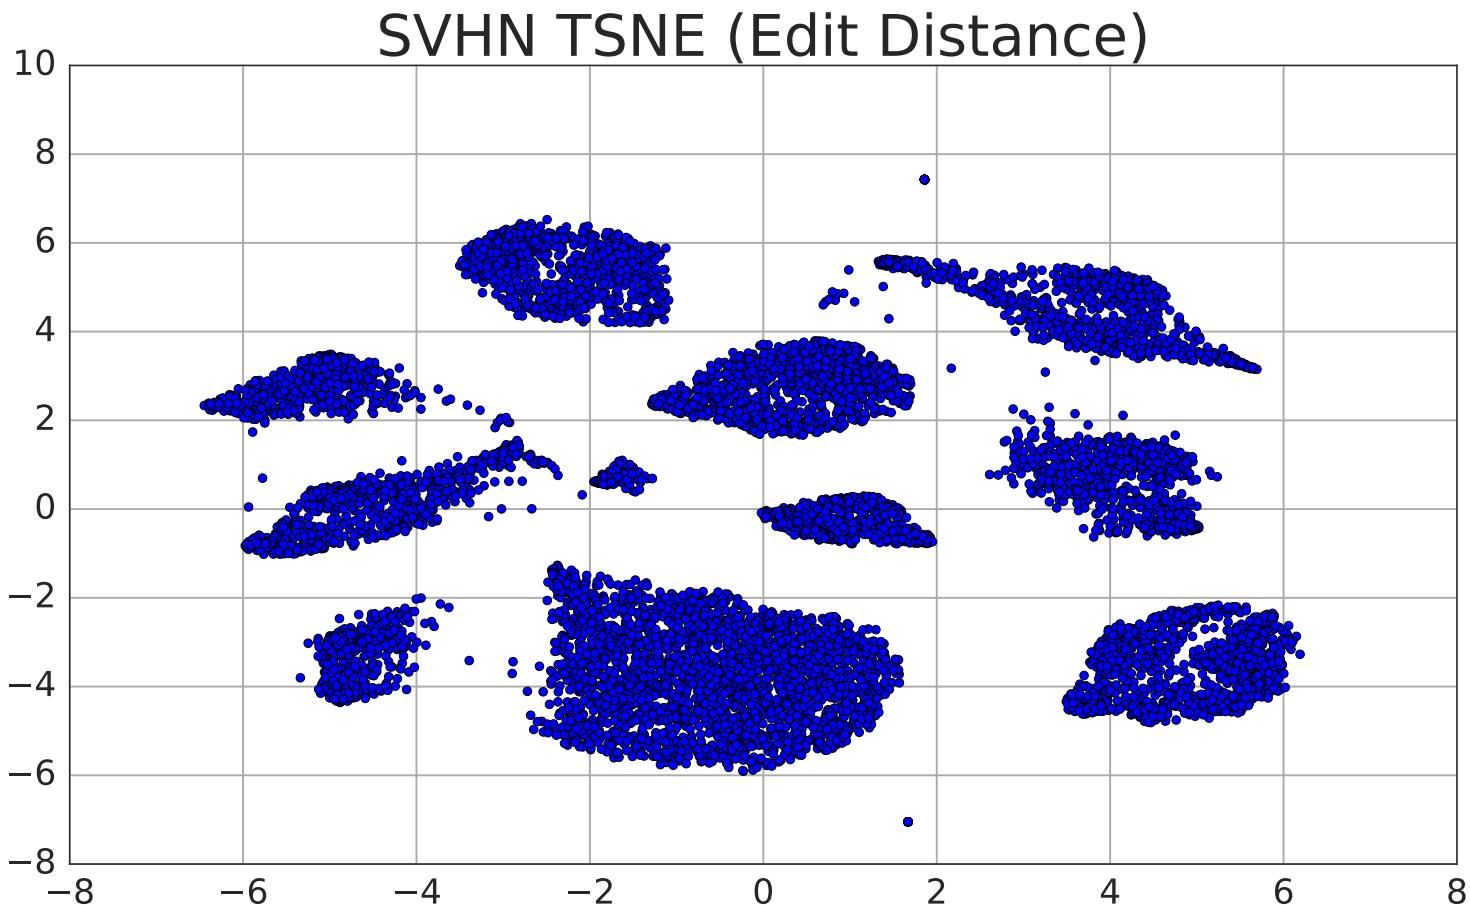
Why Does It Work?



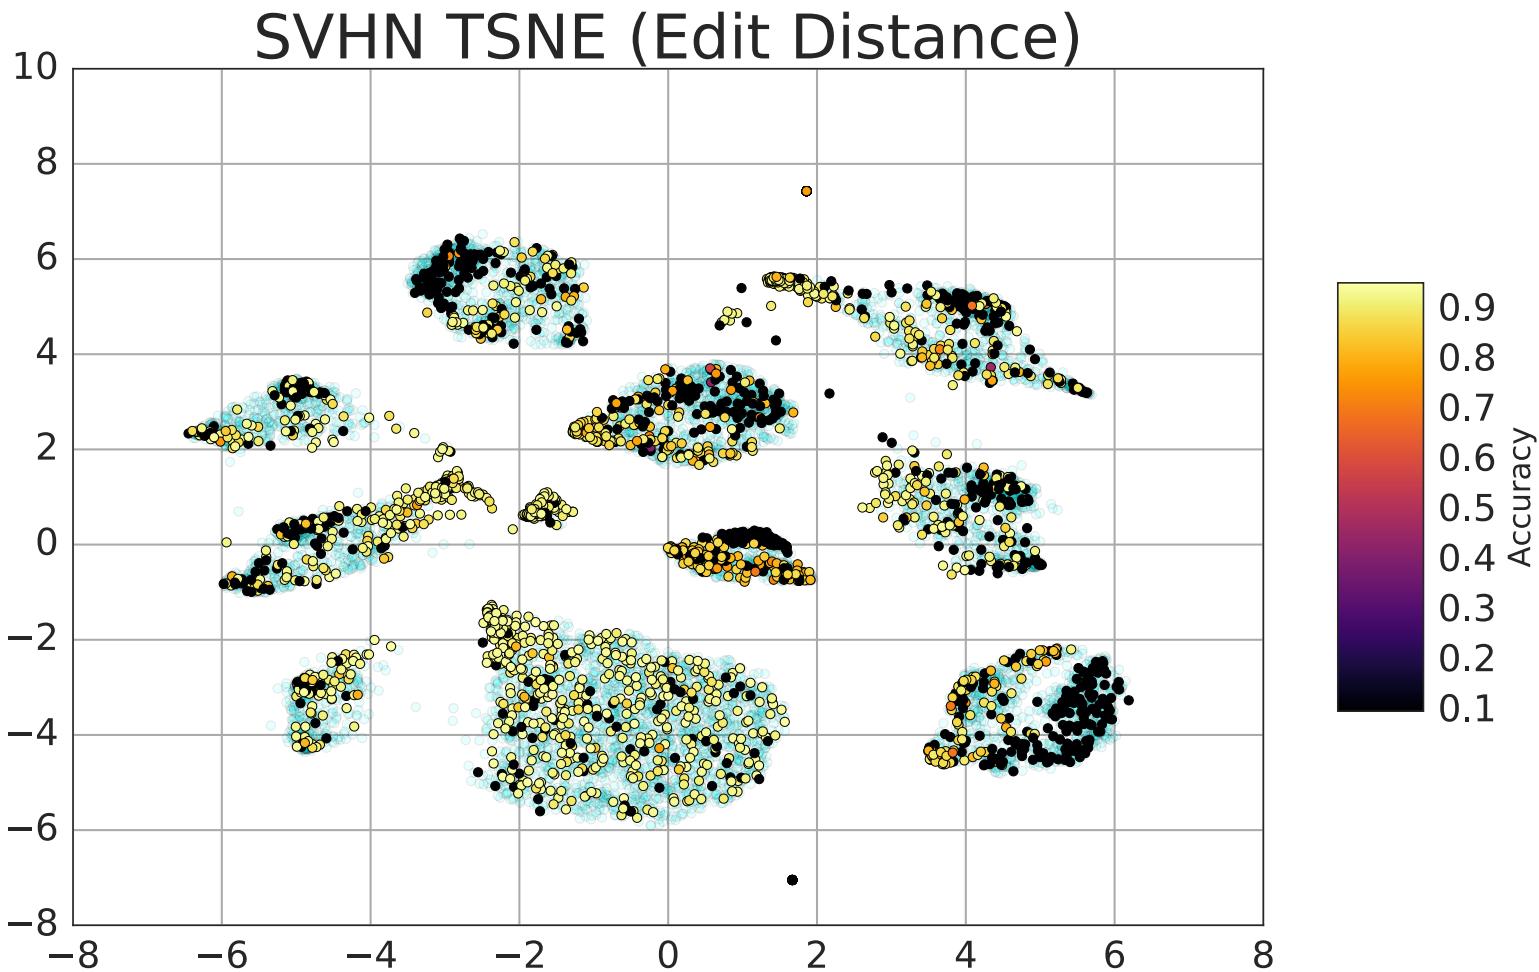
Why Does It Work?



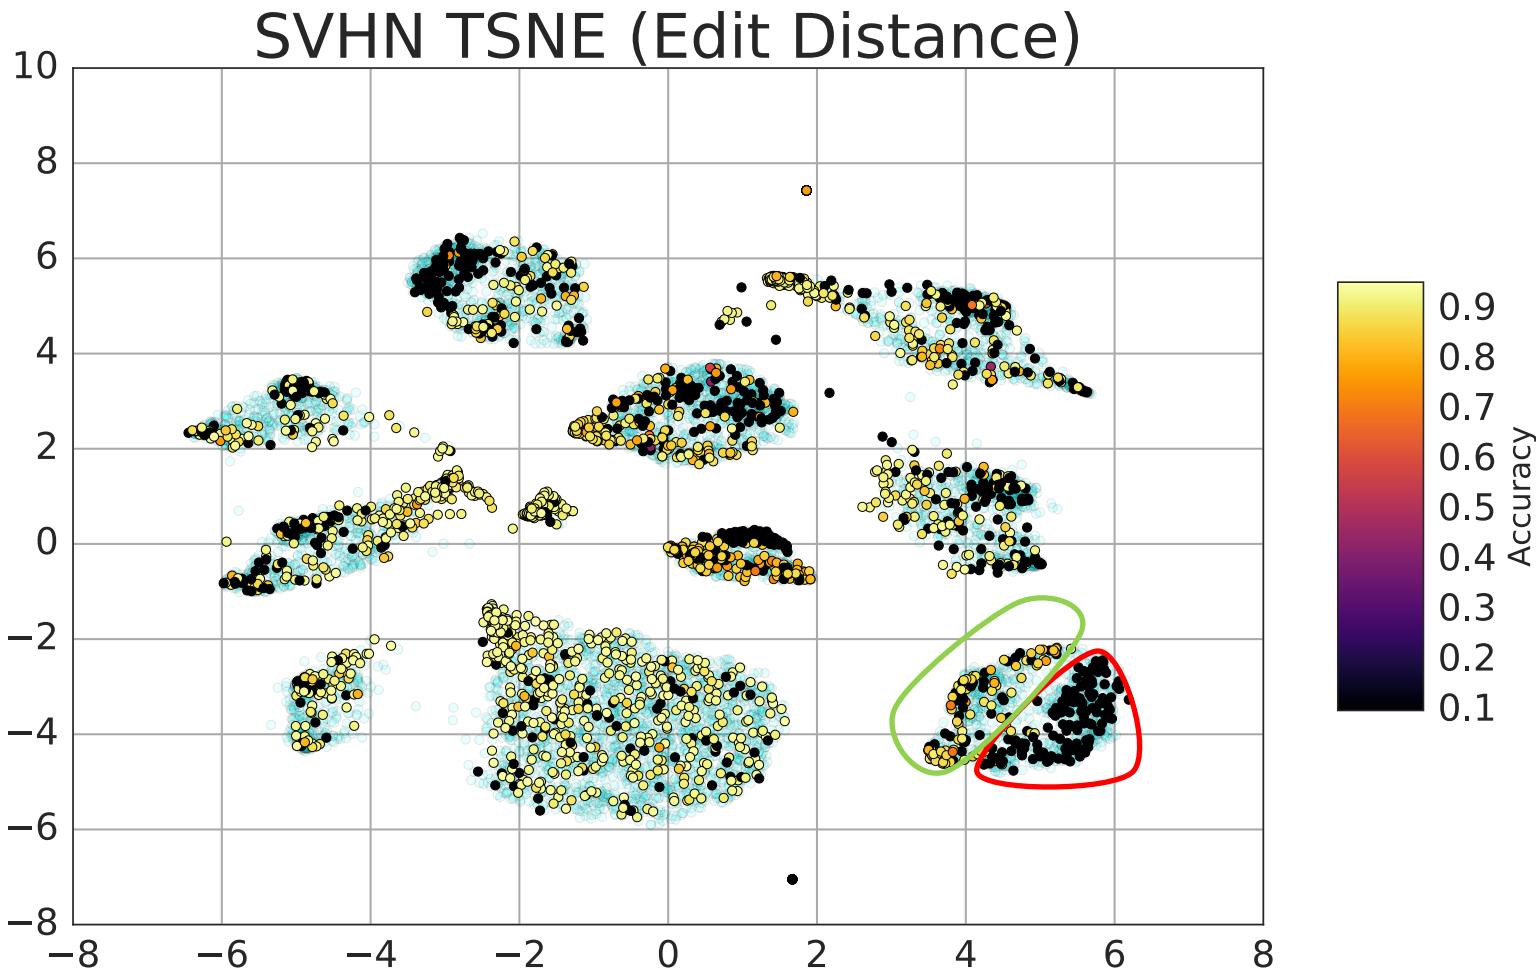
Why Does It Work?



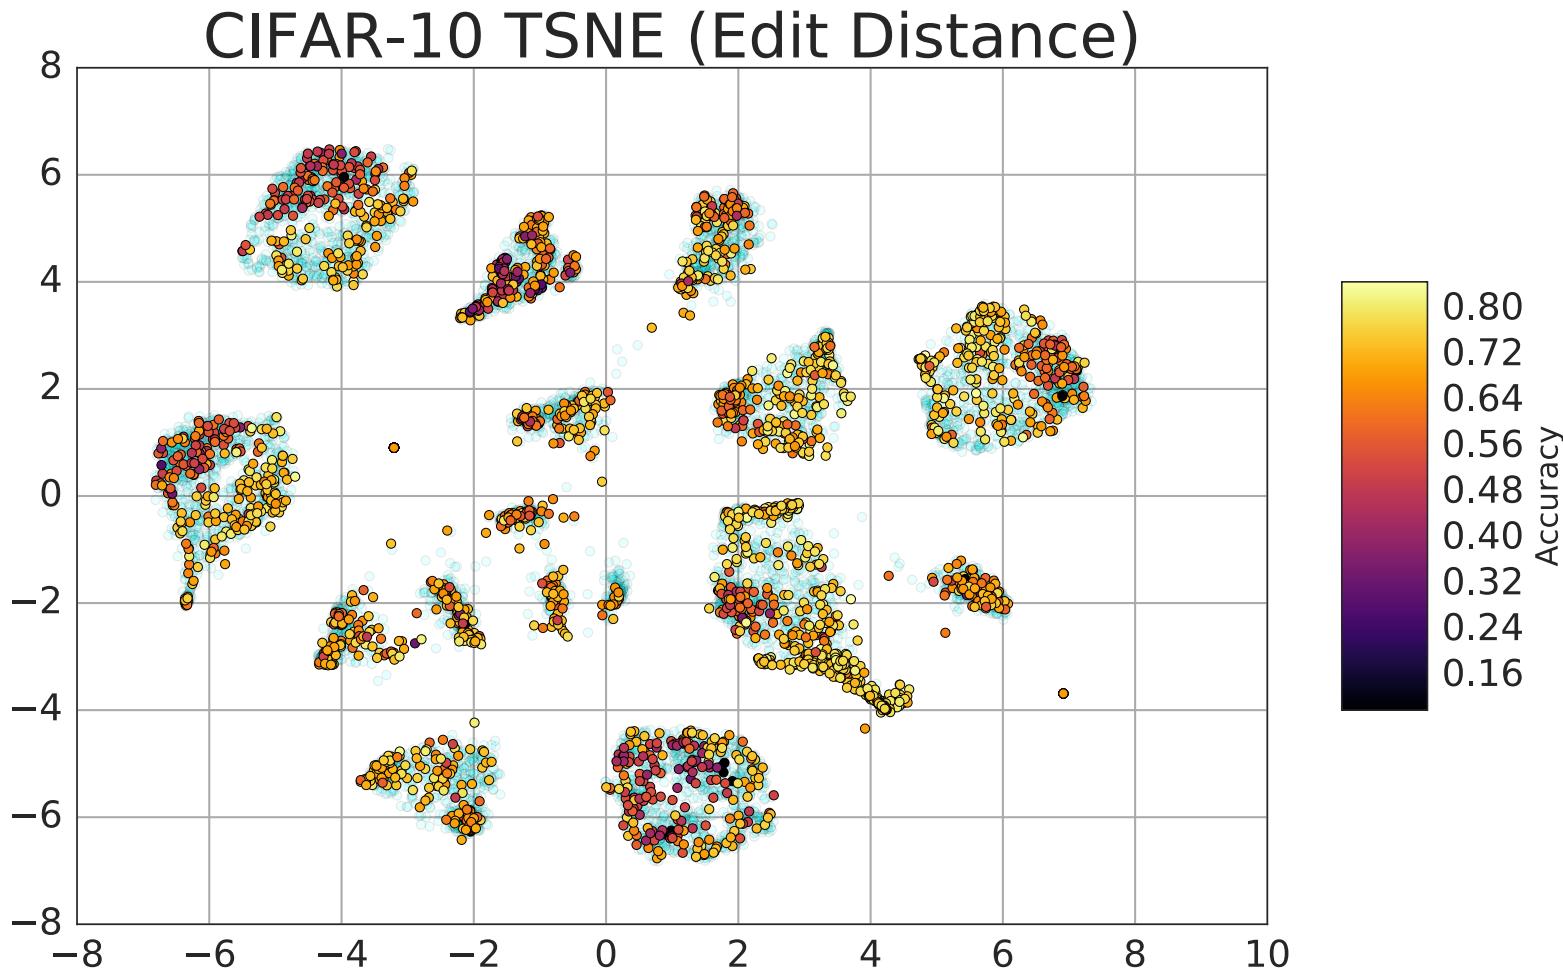
Why Does It Work?



Why Does It Work?



Why Does It Work?



Outline

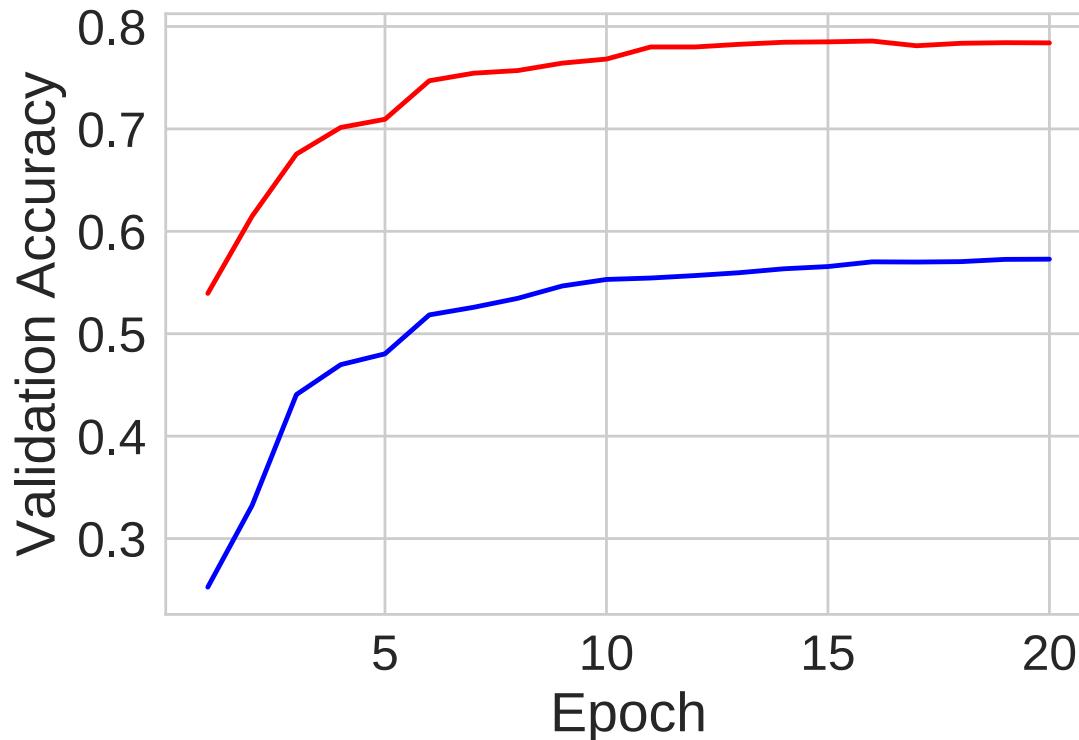
1. Reinforcement Learning Background
2. Modeling Architecture Selection as a Markov Decision Process
3. Results with Q-Learning
4. **Accelerating Architecture Selection with Simple Early Stopping Algorithms**

Meta-Modeling Comparison on CIFAR-10

| Method | Test Error on CIFAR-10 | # Samples | Estimated Computation (GPU-Days) |
|--|------------------------|-----------|----------------------------------|
| MetaQNN (Ours) | 6.92 | 2,700 | 100 |
| Neural Architecture Search (Zoph et al., 2016) | 3.65 | 12,800 | 10,000 |
| Large Scale Evolution (Real et al., 2017) | 5.4 | - | 2,600 |
| Bayesian Optimization (Snoek et al., 2012) | 9.5 | 50 | - |

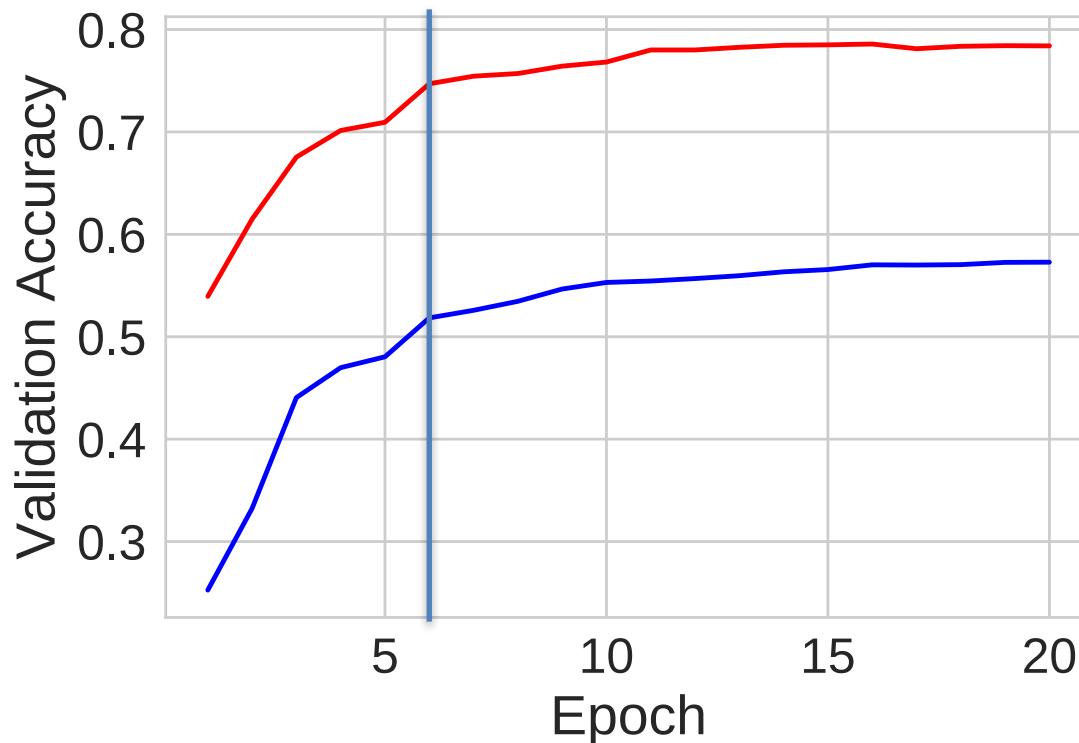
Early Stopping

- Humans are pretty good at recognizing sub-optimal training configurations



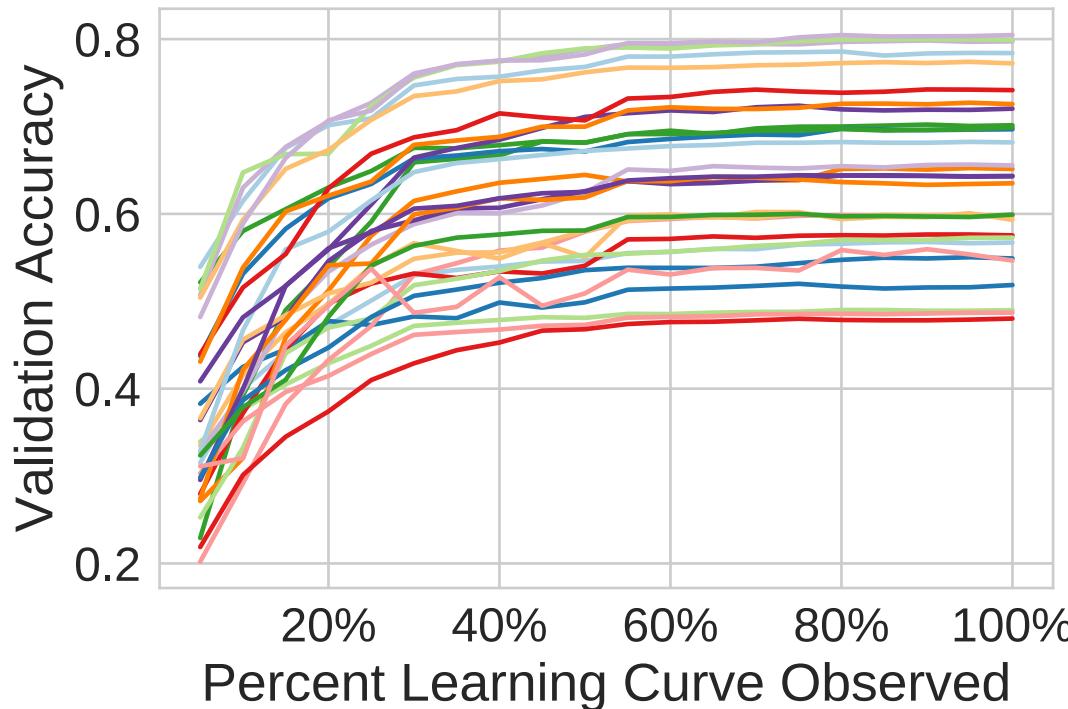
Early Stopping

- Humans are pretty good at recognizing sub-optimal training configurations



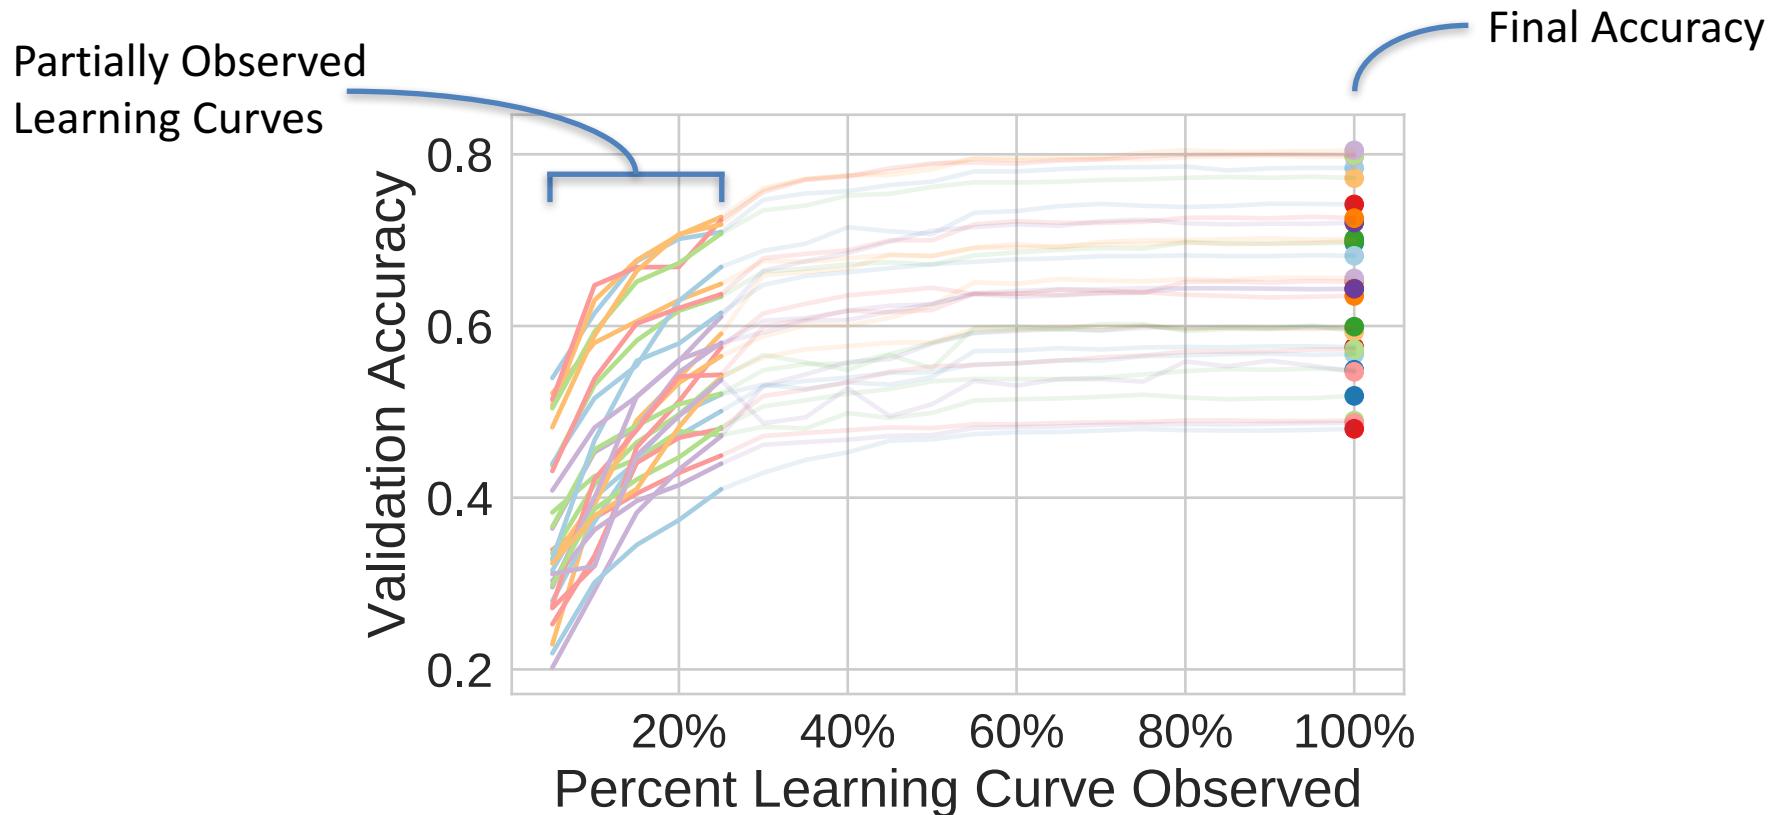
Early Stopping Using Partially Observed Learning Curves

- Use a simple model to predict final accuracy given a partially observed learning curve



Early Stopping Using Partially Observed Learning Curves

- Use a simple model to predict final accuracy given a partially observed learning curve

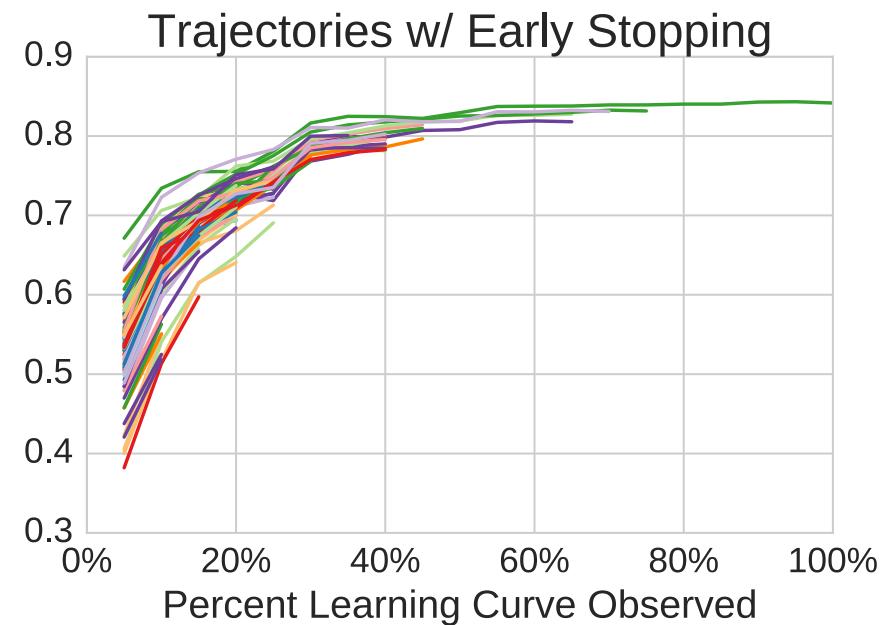
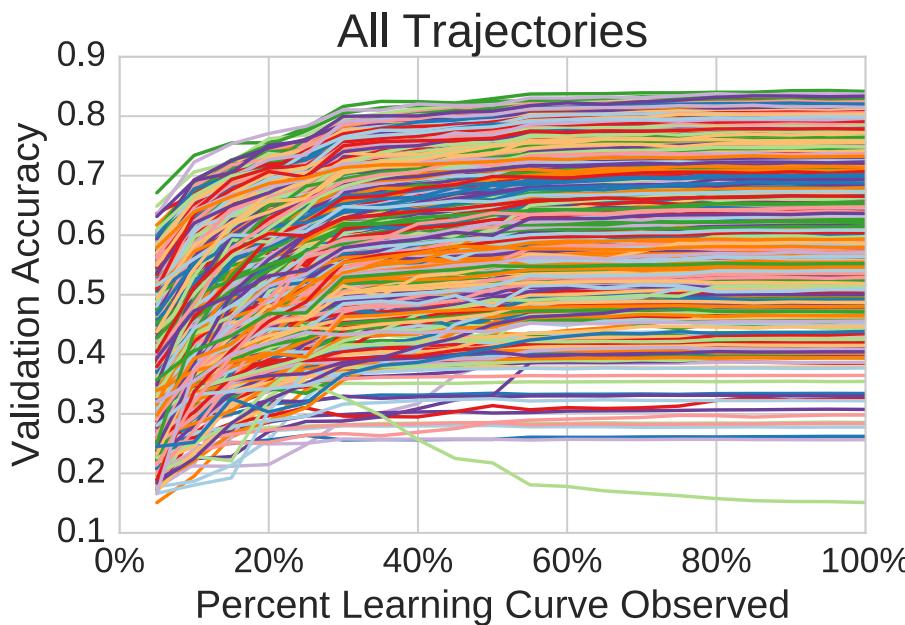


Early Stopping Using Partially Observed Learning Curves

- Use a simple model to predict final accuracy given a partially observed learning curve
- Use performance prediction to terminate sub-optimal configurations

Early Stopping Using Partially Observed Learning Curves

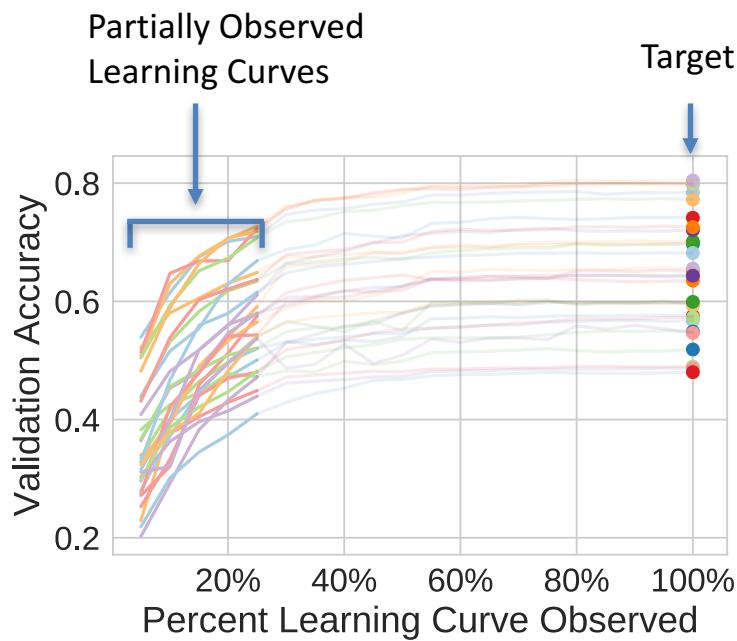
- Use a simple model to predict final accuracy given a partially observed learning curve
- Use performance prediction to terminate sub-optimal configurations



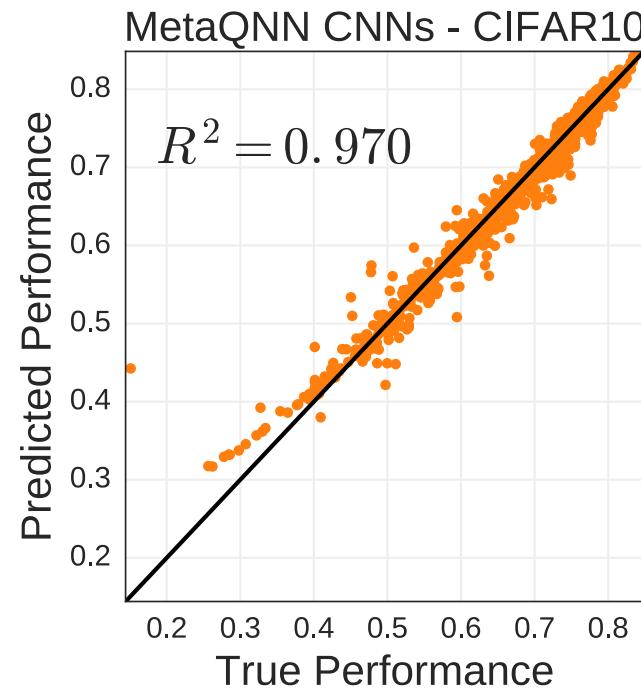
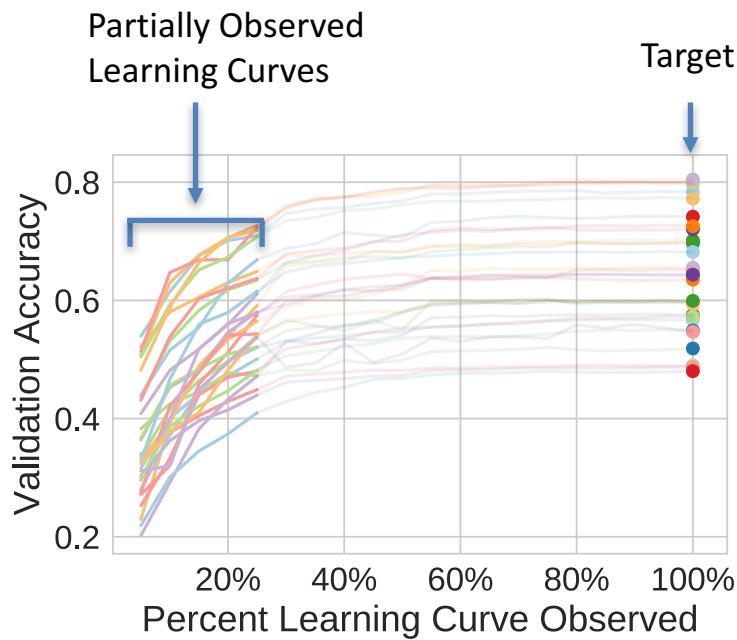
Performance Prediction Model

- Features:
 - $y_{1\dots t}$ Partially observed learning curves
 - x_f Model features, e.g. # layers, # weights, etc.
- Target
 - y_T Final Accuracy
- Works for both hyperparameter optimization and meta-modeling

Meta-Modeling Example (CIFAR-10)



Meta-Modeling Example (CIFAR-10)



- 100 training examples
- 25% learning curve observed

Experiments

- MetaQNN – Cifar10/SVHN
 - Vary Architectures

Experiments

- MetaQNN – Cifar10/SVHN
 - Vary Architectures
- Resnets – Cifar10
 - Similar search space to Neural Architecture Search

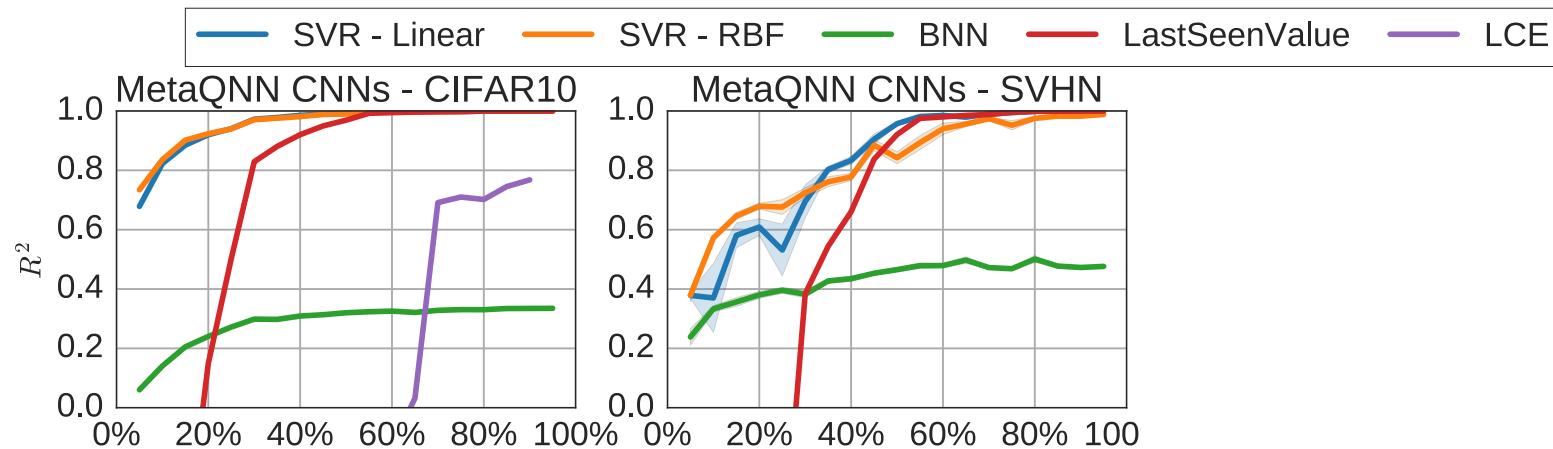
Experiments

- MetaQNN – Cifar10/SVHN
 - Vary Architectures
- Resnets – Cifar10
 - Similar search space to Neural Architecture Search
- Small Neural Network – Cifar10/SVHN
 - Vary optimization hyperparameters, e.g. learning rate, # learning rate decay steps, per layer L2 loss weight, response normalization scale and power

Experiments

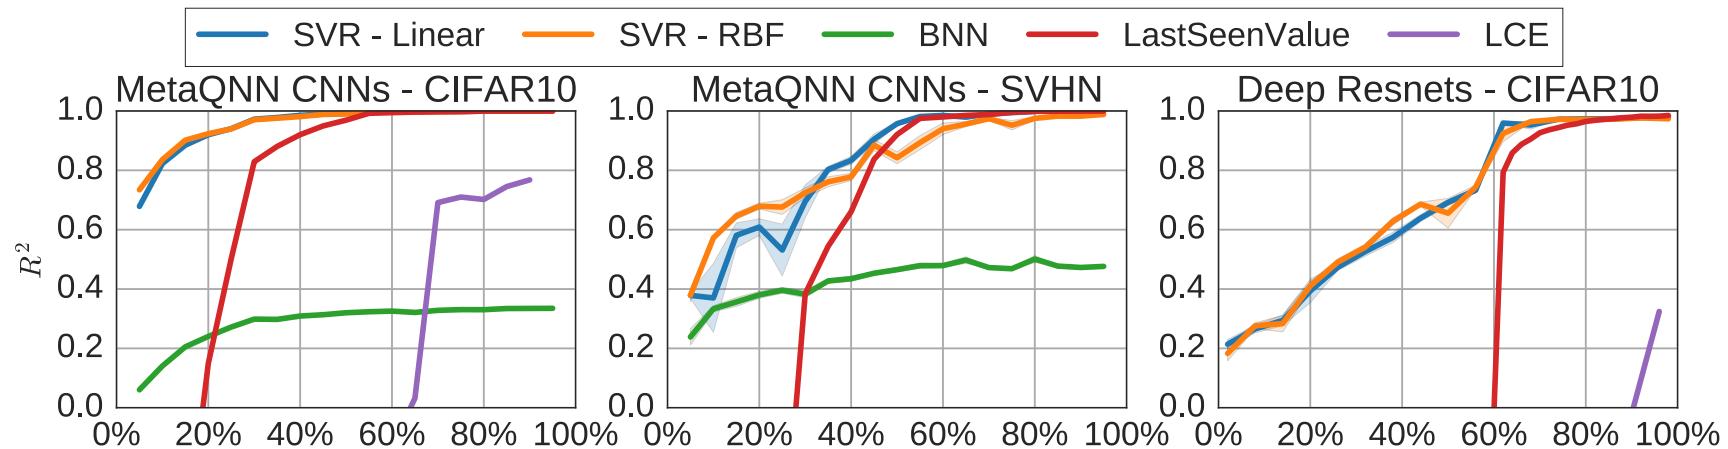
- MetaQNN – Cifar10/SVHN
 - Vary Architectures
- Resnets – Cifar10
 - Similar search space to Neural Architecture Search
- Small Neural Network – Cifar10/SVHN
 - Vary optimization hyperparameters, e.g. learning rate, # learning rate decay steps, per layer L2 loss weight, response normalization scale and power
- AlexNet – 10% ImageNet
 - Vary learning rate and # learning rate decay steps

Performance Prediction Model



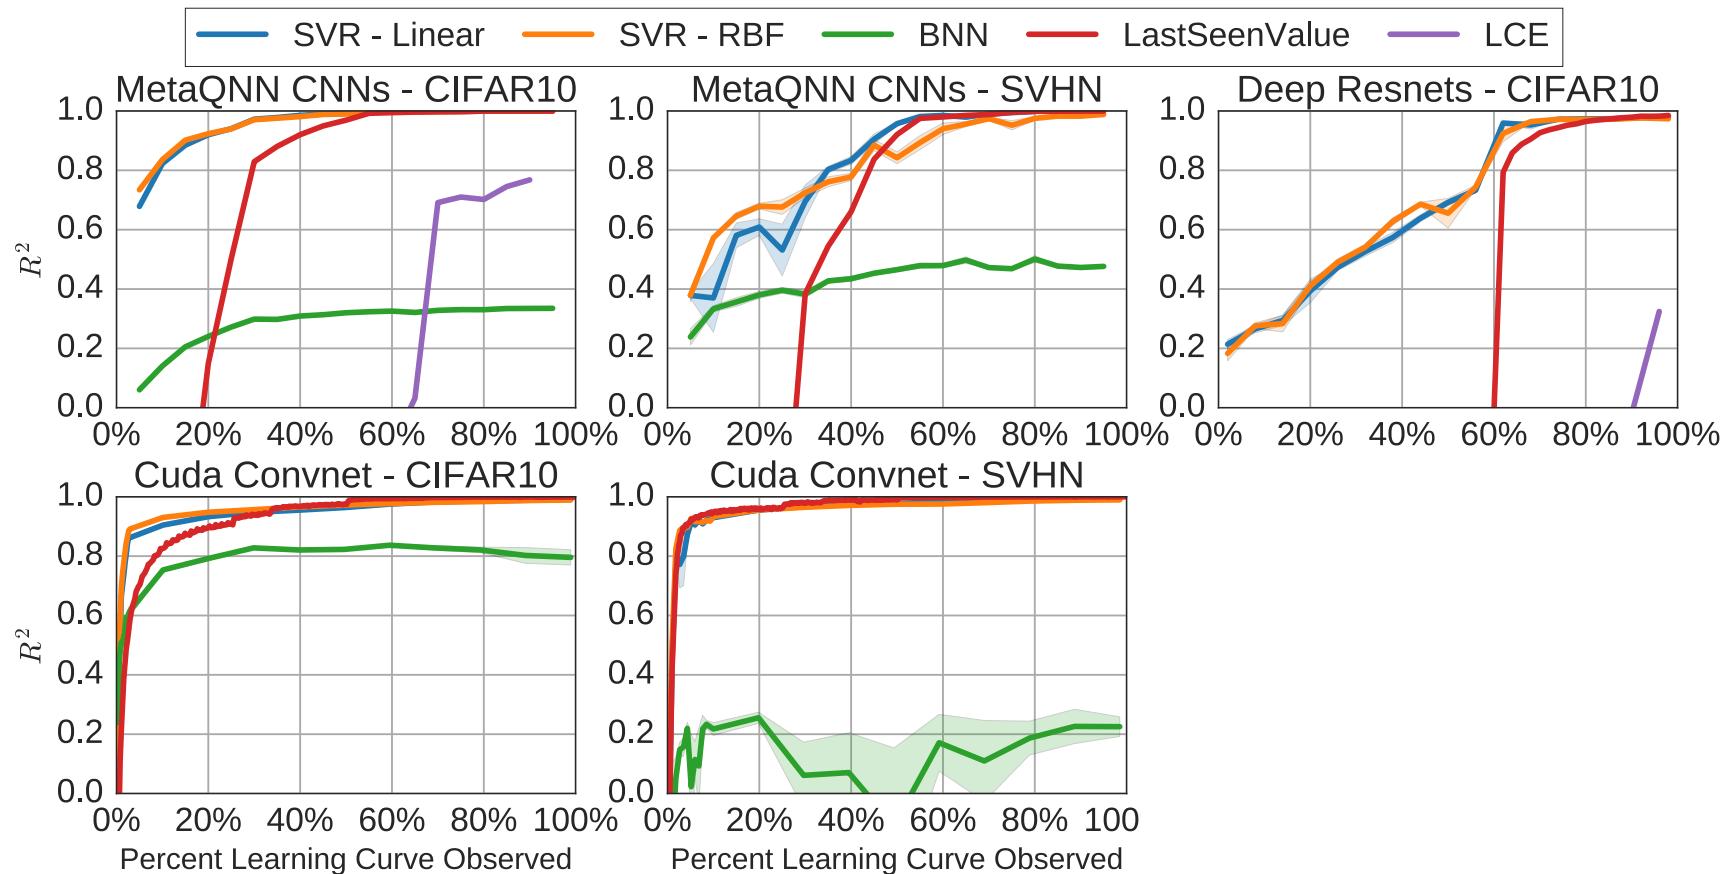
- LCE: Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves. IJCAI, 2015
- BNN: Aaron Klein, Stefan Falkner, Jost Tobias Springenberg, and Frank Hutter. Learning curve prediction with bayesian neural networks. International Conference on Learning Representations, 17, 2017.

Performance Prediction Model



- LCE: Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves. IJCAI, 2015
- BNN: Aaron Klein, Stefan Falkner, Jost Tobias Springenberg, and Frank Hutter. Learning curve prediction with bayesian neural networks. International Conference on Learning Representations, 17, 2017.

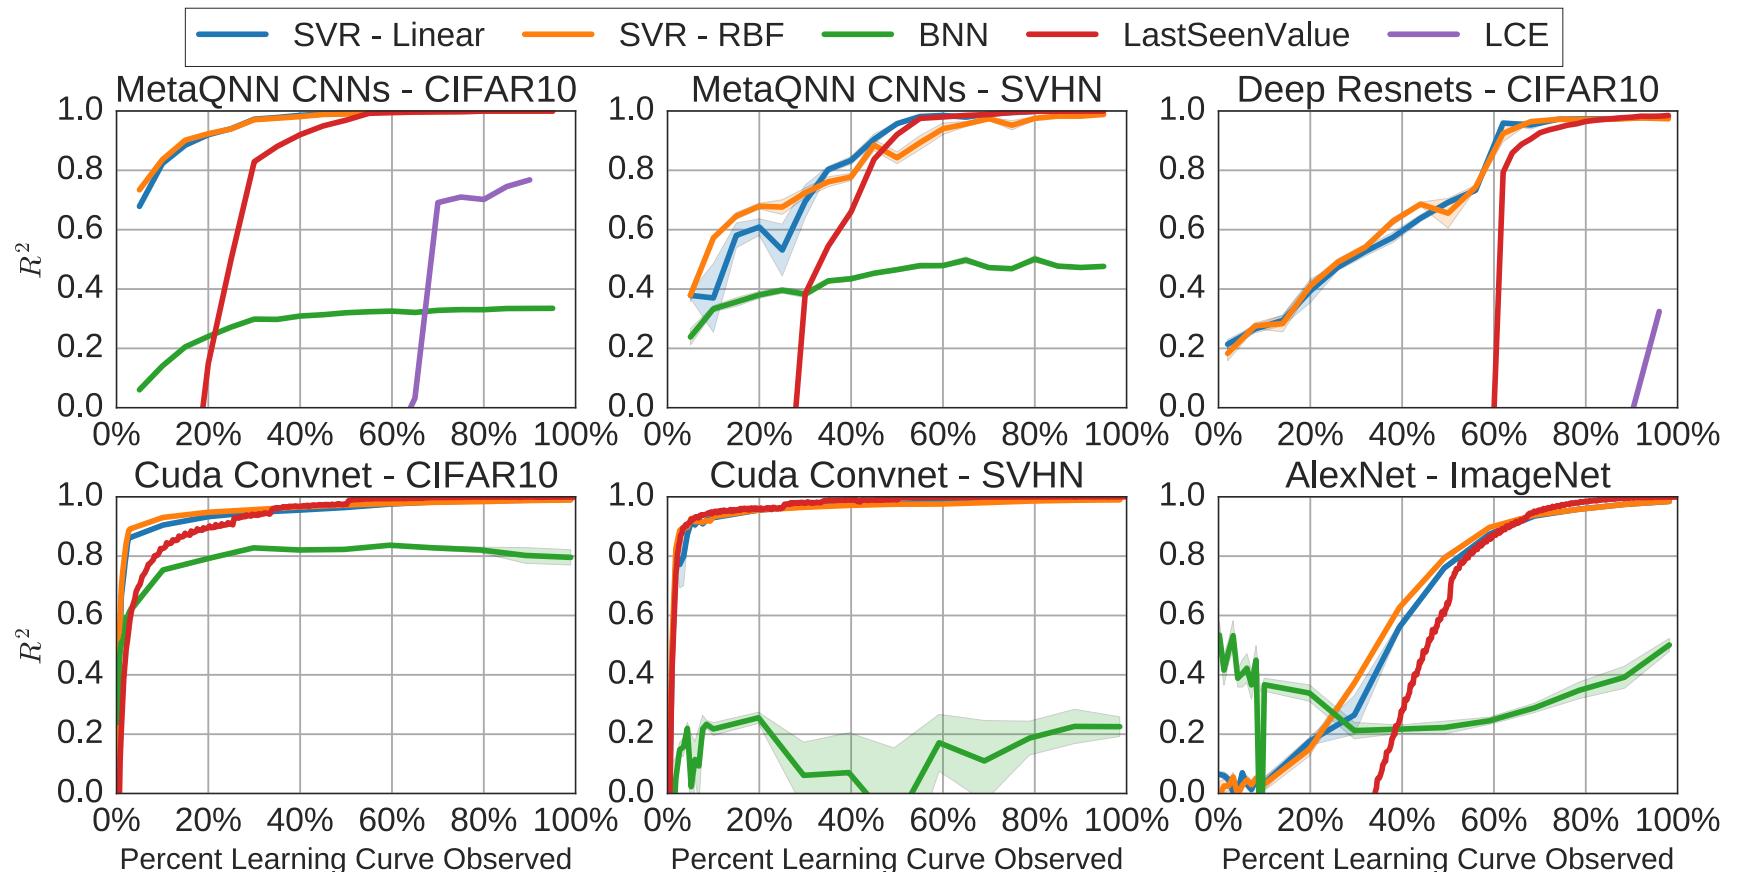
Performance Prediction Model



LCE: Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves. IJCAI, 2015

BNN: Aaron Klein, Stefan Falkner, Jost Tobias Springenberg, and Frank Hutter. Learning curve prediction with bayesian neural networks. International Conference on Learning Representations, 17, 2017.

Performance Prediction Model



- LCE: Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves. IJCAI, 2015
- BNN: Aaron Klein, Stefan Falkner, Jost Tobias Springenberg, and Frank Hutter. Learning curve prediction with bayesian neural networks. International Conference on Learning Representations, 17, 2017.

Early Stopping

1. Given performance prediction model

$$\hat{y}_T(t) = f(y_{1 \dots t}, x_f)$$

Early Stopping

1. Given performance prediction model

$$\hat{y}_T(t) = f(y_{1 \dots t}, x_f)$$

2. Assume errors are zero-mean Gaussian conditioned on t

$$\hat{y}_T(t) - y_T \sim N(0, \sigma_t)$$

Early Stopping

1. Given performance prediction model

$$\hat{y}_T(t) = f(y_{1 \dots t}, x_f)$$

2. Assume errors are zero-mean Gaussian conditioned on t

$$\hat{y}_T(t) - y_T \sim N(0, \sigma_t)$$

3. Estimate σ_t empirically from training set using LOOCV

Early Stopping

1. Given performance prediction model

$$\hat{y}_T(t) = f(y_{1 \dots t}, x_f)$$

2. Assume errors are zero-mean Gaussian conditioned on t

$$\hat{y}_T(t) - y_T \sim N(0, \sigma_t)$$

3. Estimate σ_t empirically from training set using LOOCV

4. Define probability of improvement,

$$p(\hat{y}_T(t) < y_{BEST}) = 1 - \Phi(y_{BEST}; \hat{y}_T(t), \sigma_t)$$

where $\Phi(\cdot; \mu, \sigma_t)$ is the CDF of $N(\mu, \sigma_t)$

Early Stopping

1. Given performance prediction model

$$\hat{y}_T(t) = f(y_{1 \dots t}, x_f)$$

2. Assume errors are zero-mean Gaussian conditioned on t

$$\hat{y}_T(t) - y_T \sim N(0, \sigma_t)$$

3. Estimate σ_t empirically from training set using LOOCV

4. Define probability of improvement,

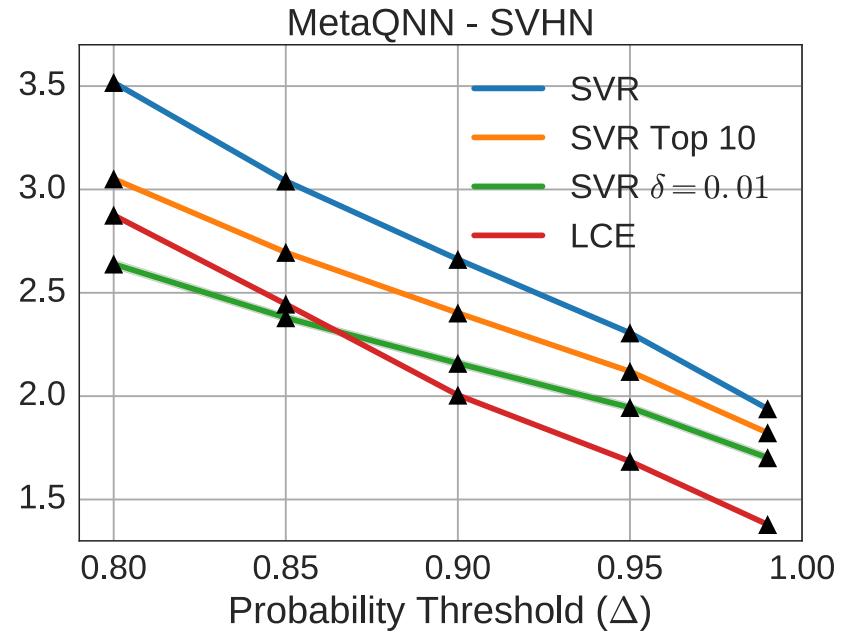
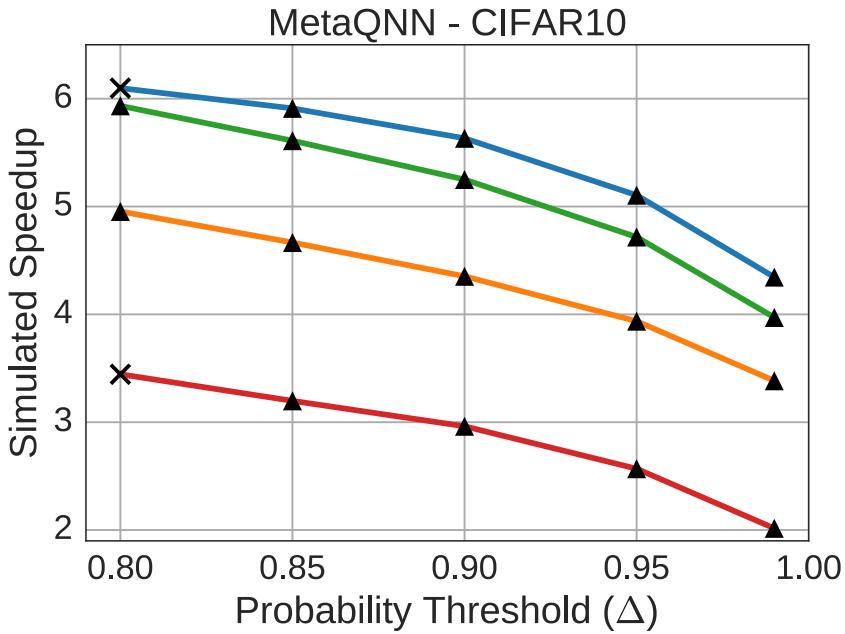
$$p(\hat{y}_T(t) < y_{BEST}) = 1 - \Phi(y_{BEST}; \hat{y}_T(t), \sigma_t)$$

where $\Phi(\cdot; \mu, \sigma_t)$ is the CDF of $N(\mu, \sigma_t)$

5. Define acceptance probability threshold Δ such that training is terminated at time-step t if

$$p(\hat{y}_T(t) < y_{BEST}) > \Delta$$

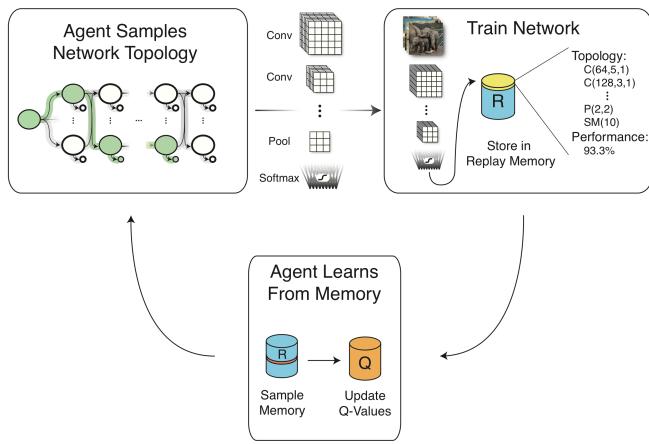
Early Stopping Results



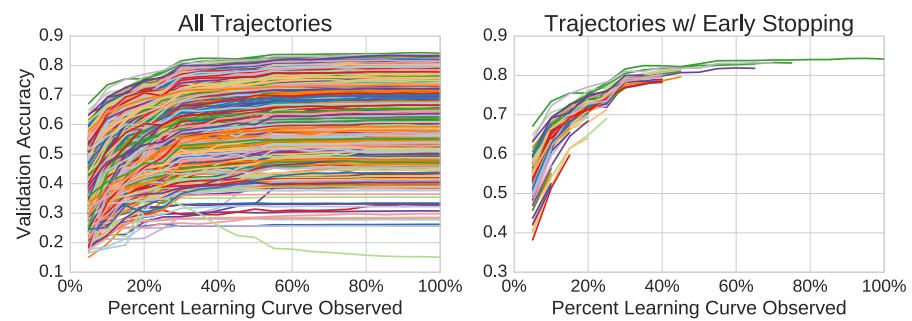
- X ~ On average does not recover best model
- ▲ ~ On average recovers best model
- δ ~ Termination rule $p(\hat{y}_T(t) < y_{BEST} - \delta) > \Delta$
- Top 10 ~ Termination rule $p(\hat{y}_T(t) < y_{10^{th} BEST}) > \Delta$

Summary

Designing neural network architectures using reinforcement learning [1]



Practical Neural Network Performance Prediction for Early Stopping [2]



Contact: bowen@mit.edu

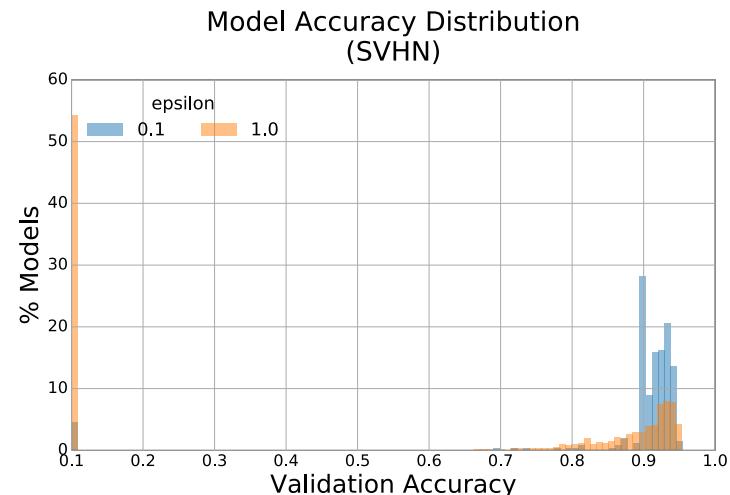
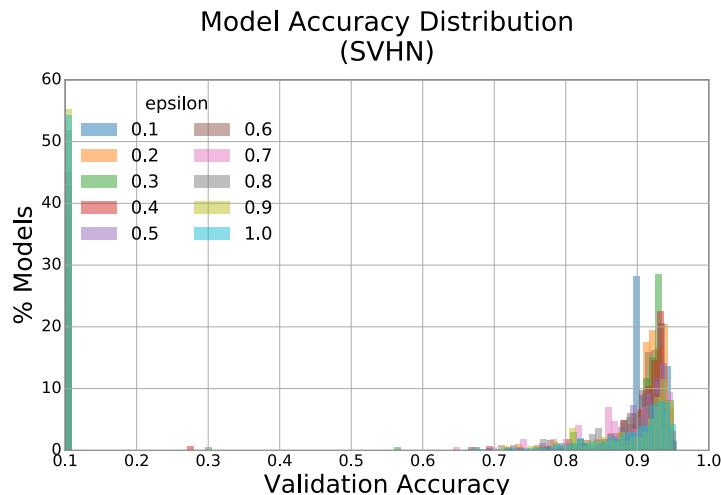
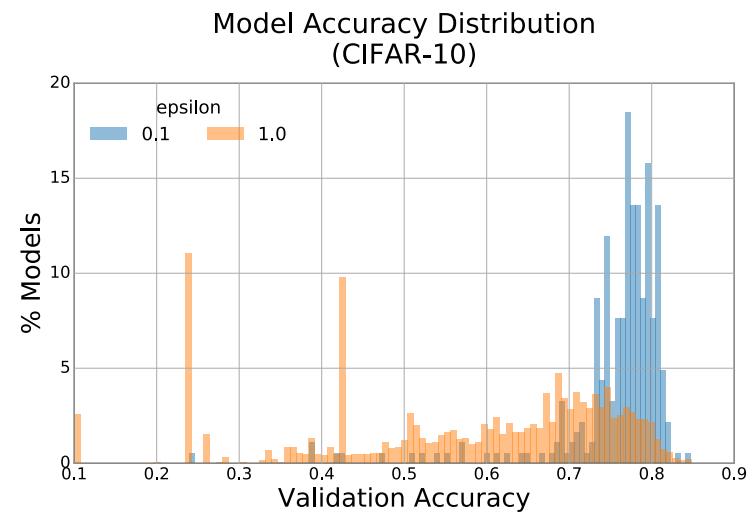
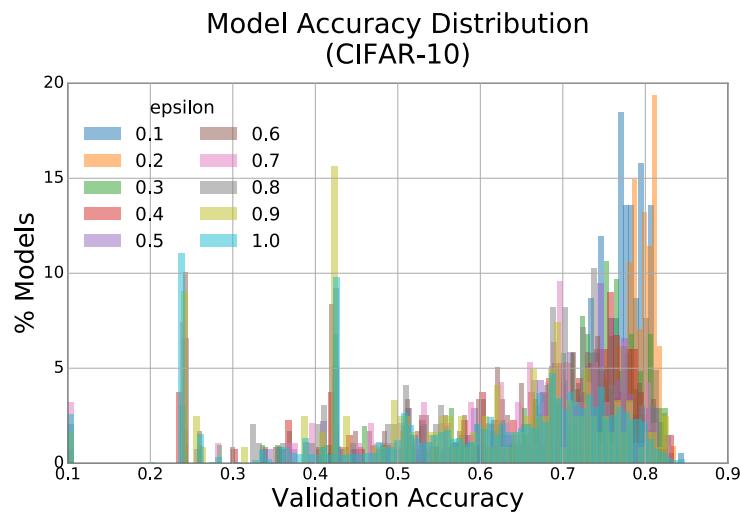
Slides: bowenbaker.github.io (check back later today)

MetaQNN Code: Released by end of week

1. Bowen Baker, Otkrist Gupta, Nikhil Naik, and Ramesh Raskar. "Designing neural network architectures using reinforcement learning." *International Conference on Learning Representations*, 2017.
2. Bowen Baker*, Otkrist Gupta*, Ramesh Raskar, and Nikhil Naik. "Practical Neural Network Performance Prediction for Early Stopping." *Under Submission*, 2017.

Appendix

Exploration Distributions

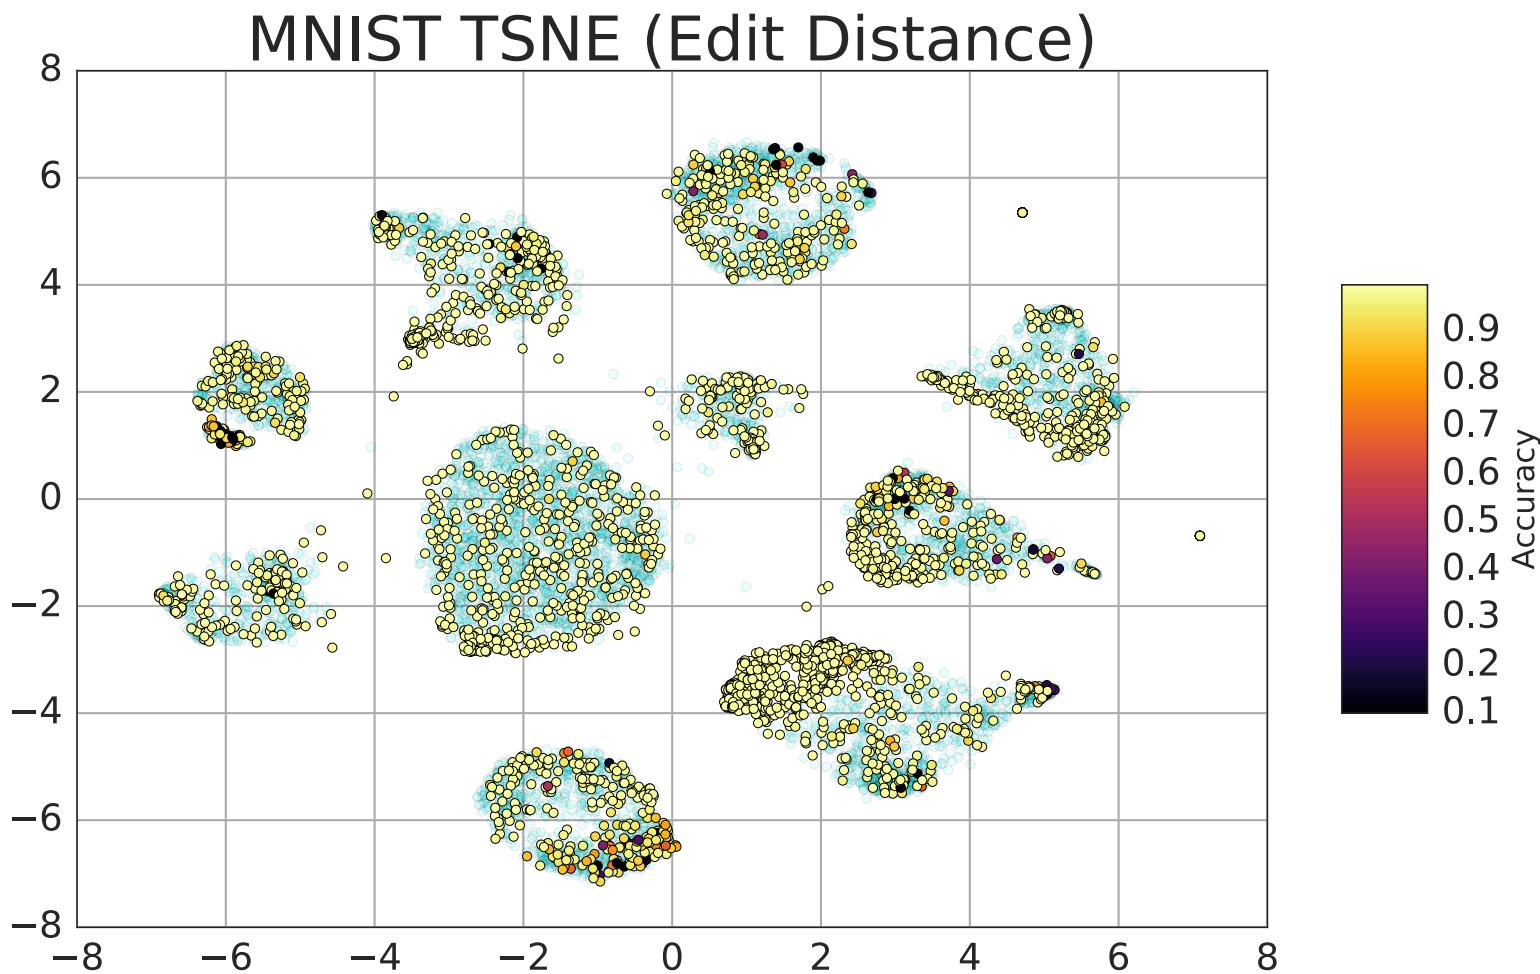


Transferability

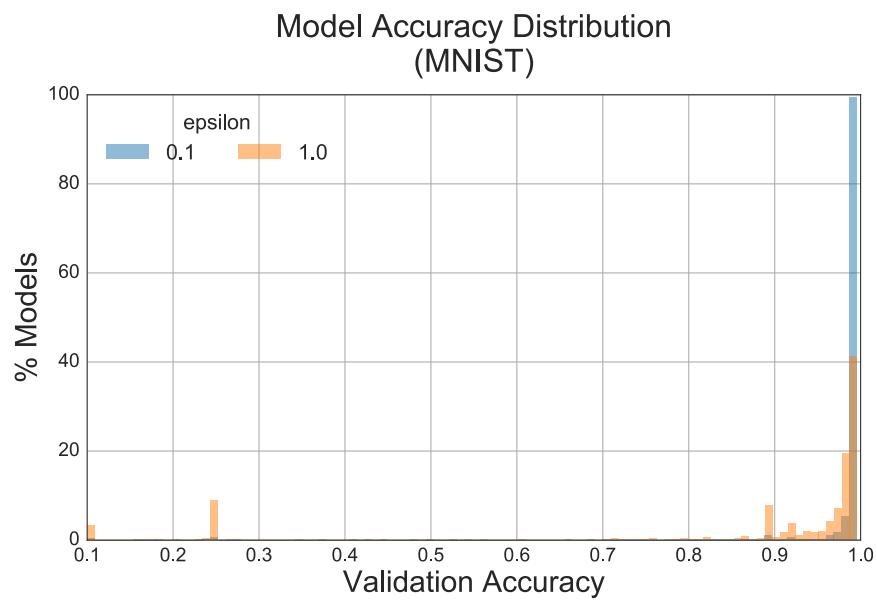
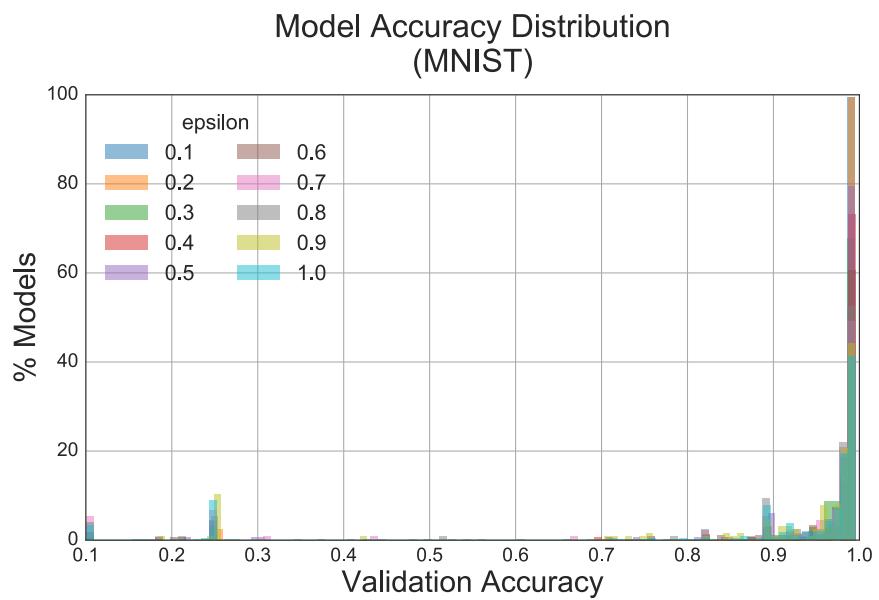
- Top model found in CIFAR-10 experiment trained for other tasks

| Dataset | CIFAR-100 | SVHN | MNIST |
|-----------------------|------------------------------|-------------------------|-------------------------|
| Training from scratch | 27.14 | 2.48 | 0.80 |
| Finetuning | 34.93 | 4.00 | 0.81 |
| State-of-the-art | 24.28 (Clevert et al., 2015) | 1.69 (Lee et al., 2016) | 0.31 (Lee et al., 2016) |

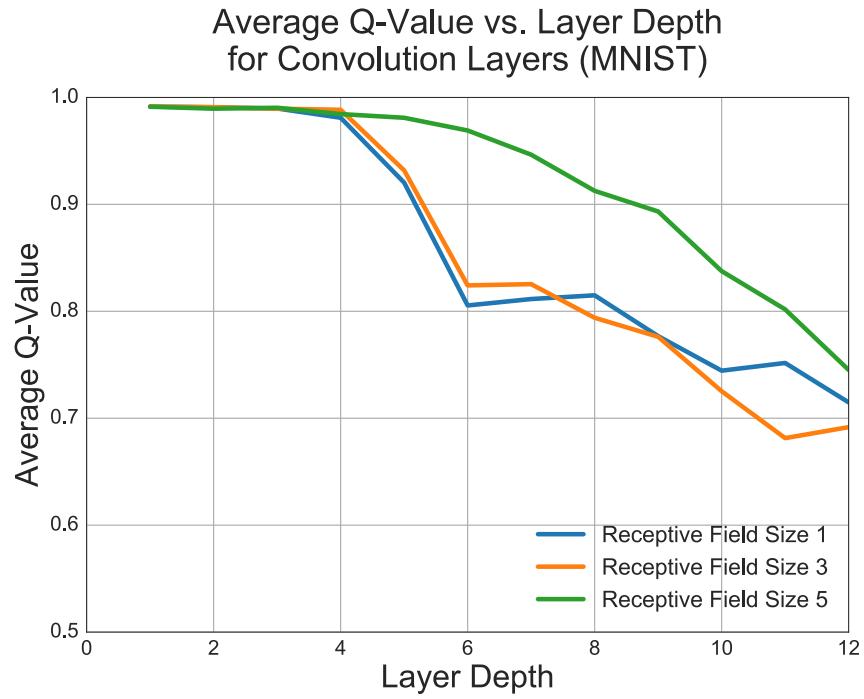
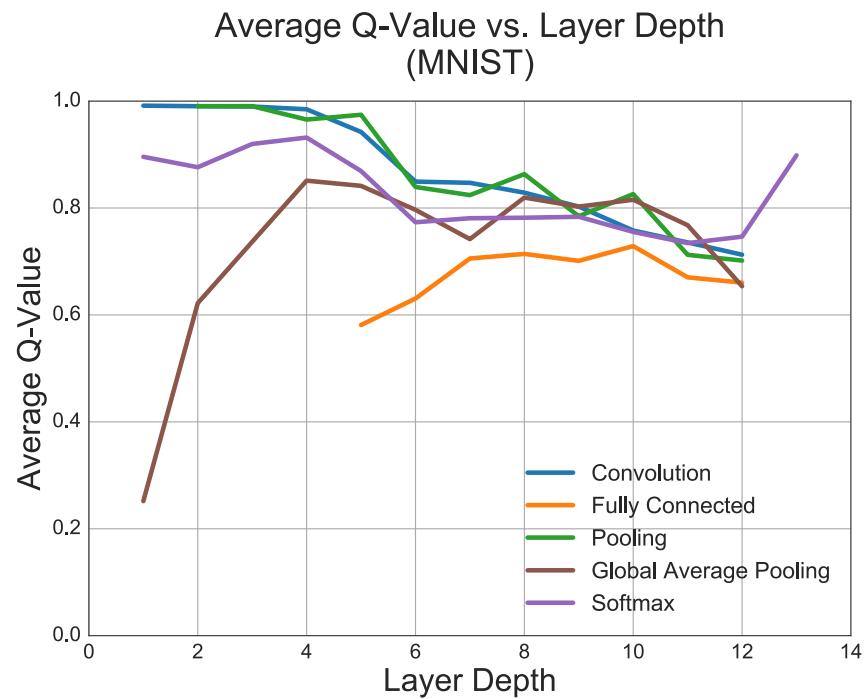
MNIST t-SNE



MNIST Exploration Distribution

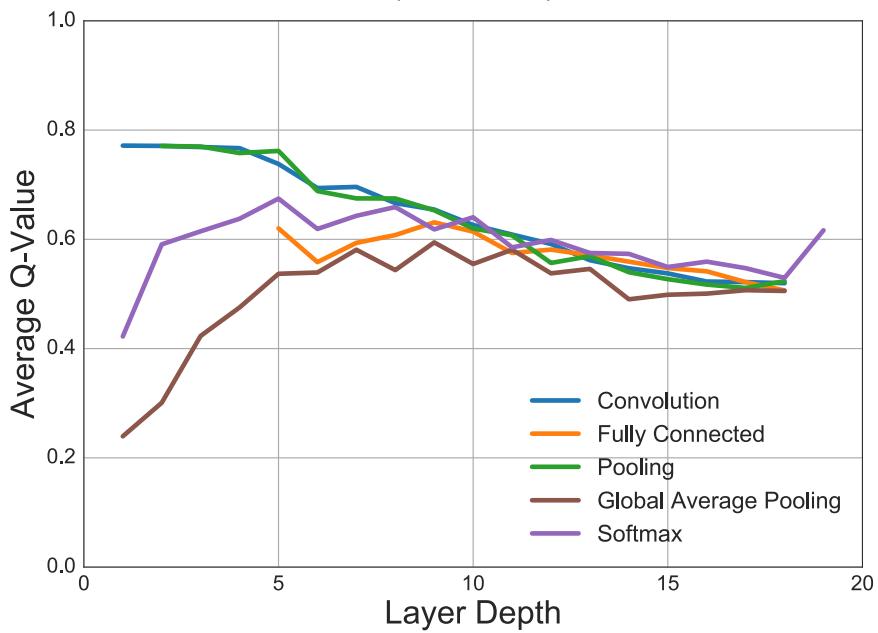


MNIST Q-Value Analysis

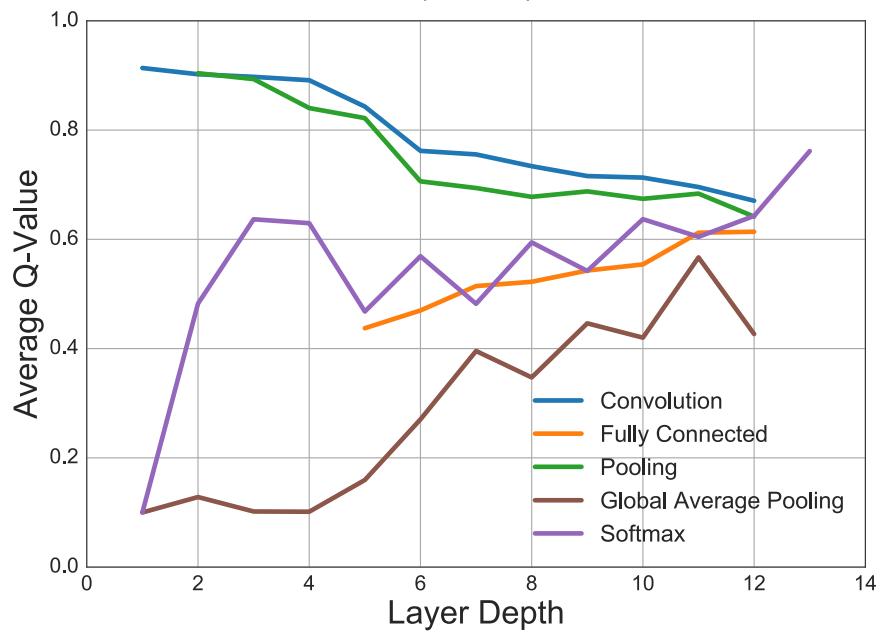


Q-Value Analysis

Average Q-Value vs. Layer Depth
(CIFAR10)



Average Q-Value vs. Layer Depth
(SVHN)



Top Models (CIFAR-10)

| Model Architecture | Test Error (%) | # Params (10^6) |
|---|----------------|---------------------|
| [C(512,5,1), C(256,3,1), C(256,5,1), C(256,3,1), P(5,3), C(512,3,1), C(512,5,1), P(2,2), SM(10)] | 6.92 | 11.18 |
| [C(128,1,1), C(512,3,1), C(64,1,1), C(128,3,1), P(2,2), C(256,3,1), P(2,2), C(512,3,1), P(3,2), SM(10)] | 8.78 | 2.17 |
| [C(128,3,1), C(128,1,1), C(512,5,1), P(2,2), C(128,3,1), P(2,2), C(64,3,1), C(64,5,1), SM(10)] | 8.88 | 2.42 |
| [C(256,3,1), C(256,3,1), P(5,3), C(256,1,1), C(128,3,1), P(2,2), C(128,3,1), SM(10)] | 9.24 | 1.10 |
| [C(128,5,1), C(512,3,1), P(2,2), C(128,1,1), C(128,5,1), P(3,2), C(512,3,1), SM(10)] | 11.63 | 1.66 |

Top Models (SVHN)

| Model Architecture | Test Error (%) | # Params (10^6) |
|--|----------------|---------------------|
| [C(128,3,1), P(2,2), C(64,5,1), C(512,5,1), C(256,3,1), C(512,3,1), P(2,2), C(512,3,1), C(256,5,1), C(256,3,1), C(128,5,1), C(64,3,1), SM(10)] | 2.24 | 9.81 |
| [C(128,1,1), C(256,5,1), C(128,5,1), P(2,2), C(256,5,1), C(256,1,1), C(256,3,1), C(256,3,1), C(256,5,1), C(512,5,1), C(256,3,1), C(128,3,1), SM(10)] | 2.28 | 10.38 |
| [C(128,5,1), C(128,3,1), C(64,5,1), P(5,3), C(128,3,1), C(512,5,1), C(256,5,1), C(128,5,1), C(128,5,1), C(128,3,1), SM(10)] | 2.32 | 6.83 |
| [C(128,1,1), C(256,5,1), C(128,5,1), C(256,3,1), C(256,5,1), P(2,2), C(128,1,1), C(512,3,1), C(256,5,1), P(2,2), C(64,5,1), C(64,1,1), SM(10)] | 2.35 | 6.99 |
| [C(128,1,1), C(256,5,1), C(128,5,1), C(256,5,1), C(256,5,1), C(256,1,1), P(3,2), C(128,1,1), C(256,5,1), C(512,5,1), C(256,3,1), C(128,3,1), SM(10)] | 2.36 | 10.05 |

Top Models (MNIST)

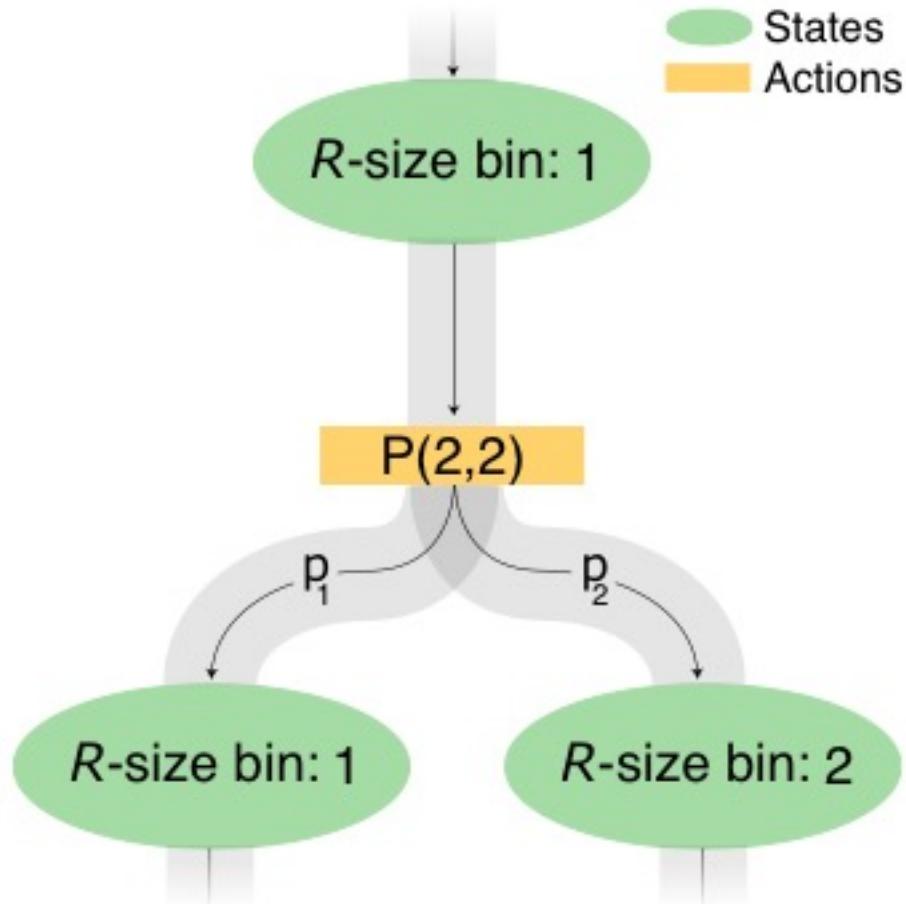
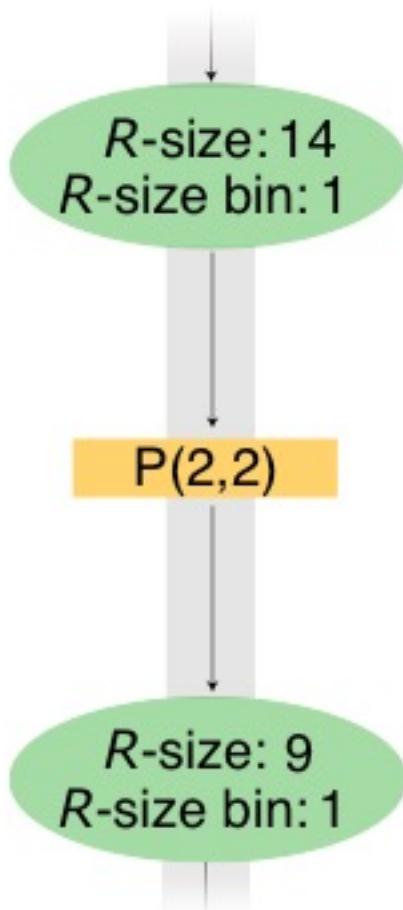
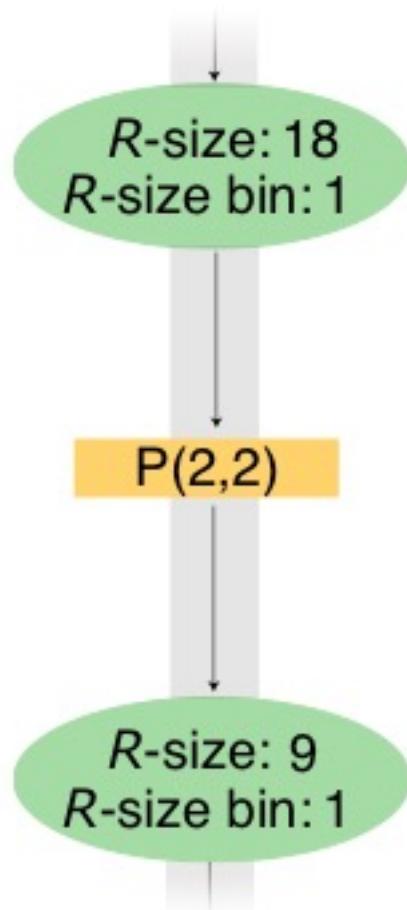
| Model Architecture | Test Error (%) | # Params (10^6) |
|--|----------------|---------------------|
| [C(64,1,1), C(256,3,1), P(2,2), C(512,3,1), C(256,1,1), P(5,3), C(256,3,1), C(512,3,1), FC(512), SM(10)] | 0.35 | 5.59 |
| [C(128,3,1), C(64,1,1), C(64,3,1), C(64,5,1), P(2,2), C(128,3,1), P(3,2), C(512,3,1), FC(512), FC(128), SM(10)] | 0.38 | 7.43 |
| [C(512,1,1), C(128,3,1), C(128,5,1), C(64,1,1), C(256,5,1), C(64,1,1), P(5,3), C(512,1,1), C(512,3,1), C(256,3,1), C(256,5,1), C(256,5,1), SM(10)] | 0.40 | 8.28 |
| [C(64,3,1), C(128,3,1), C(512,1,1), C(256,1,1), C(256,5,1), C(128,3,1), P(5,3), C(512,1,1), C(512,3,1), C(128,5,1), SM(10)] | 0.41 | 6.27 |
| [C(64,3,1), C(128,1,1), P(2,2), C(256,3,1), C(128,5,1), C(64,1,1), C(512,5,1), C(128,5,1), C(64,1,1), C(512,5,1), C(256,5,1), C(64,5,1), SM(10)] | 0.43 | 8.10 |

Top Model Cifar-10 (Updated Results)

```
[C(64,3,1), C(256,3,1), D(1,9), C(512,3,1), C(64,3,1),  
D(2,9), C(128,5,1), P(2,2), D(3,9), C(512,5,1), P(2,2),  
D(4,9), C(128,5,1), C(256,5,1), D(5,9), C(512,3,1),  
C(64,5,1), D(6,9), P(2,2), C(512,1,1), D(7,9), FC(128),  
D(8,9), SM(10)]
```

Representation Size

States
Actions



Q-Learning

$Q^*(s, u)$ -- Denotes the expected reward when following an optimal policy after taking action u at state s

Q-Learning

$$Q^*(s_i, u) = \mathbb{E} \left[r + \gamma \max_{u' \in \mathcal{U}(s_j)} Q^*(s_j, u') \right]$$

γ -- Discount Factor

r -- Reward received from
the (s_i, u, s_j) transition

Q-Learning

$$Q^*(s_i, u) = \mathbb{E} \left[r + \gamma \max_{u' \in \mathcal{U}(s_j)} Q^*(s_j, u') \right]$$

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha \left[r_t + \gamma \max_{u' \in \mathcal{U}(s_j)} Q_t(s_j, u') \right]$$